Complete Cross-Cloud SOC Auto-Remediation System

Production-Ready Security Automation | Microsoft Sentinel + AWS Lambda + API Gateway + Amazon SES + DynamoDB

Microsoft Sentinel | Azure Logic Apps | KQL | AWS Lambda | API Gateway | IAM | SES | DynamoDB | Python 3.12 | OIDC | CloudTrail

Musa Olalekan Cloud Sec | Workspace: soclawlab | Analytics Rules: 5 | Attacks Simulated: 5 | Incidents: 2 | Verdict: CRITICAL | Response < 30 sec

Complete Documentation Index (All 28 Phases Preserved - Grade A Master Documentation)

- 1. Executive Summary & Project Overview
- 2. Complete Architecture & Flow Diagram
- 3. Prerequisites & Environment Setup

- 4. PHASE 1-3: Log Analytics Workspace & Sentinel Enablement
- 5. PHASE 4-5: Azure Activity & Entra ID Log Integration
- 6. PHASE 6: Wait for Data Ingestion
- 7. PHASE 7: Verify Logs & First KQL Queries
- 8. PHASE 8: Analyze Logs for Security Insights
- 9. PHASE 9: Workbooks & Visualizations
- 10. PHASE 10: 5 Analytics Rules (Full KQL + MITRE ATT&CK Mapping)
- 11. PHASE 11: Attack Simulations (5 Scenarios - Full Details + Business Impact)
- 12. PHASE 12: Advanced KQL Hunting Queries (14+ Queries)
- 13. PHASE 13: Forensic Investigation & Incident Report (Root Cause Analysis)
- 14. PHASE 14: Logic App Designer & Email Automation (I Built This FIRST)
- 15. PHASE 15: Azure AD Auto-Remediation (Privilege Escalation Containment)
- 16. PHASE 16: Connecting AWS to Azure Sentinel (OIDC Integration)
- 17. PHASE 17: Amazon SES Email Verification
- 18. PHASE 18: DynamoDB Idempotency Table
- 19. PHASE 19: Lambda Function (Complete Python Code)
- 20. PHASE 20: API Gateway Secure Entry Point
- 21. PHASE 21: Logic App Cross-Cloud Orchestration
- 22. PHASE 22: IAM Least Privilege Policy
- 23. PHASE 23: Testing & Validation Results
- 24. PHASE 24: Automation Rule & Circuit Breaker (Fail-Safe Mechanisms)
- 25. Interview Talking Points & Answers
- 26. Portfolio Checklist & CV Bullets
- 27. Appendix: All Code Listings (Full)
- 28. Lessons Learned & Future Enhancements

1. Executive Summary & Project Overview

This report documents the complete implementation of a **production-grade cross-cloud SOC auto-remediation system.** The solution bridges Microsoft Azure Sentinel and Amazon Web Services (AWS) to deliver automated threat detection, real-time alerting, and immediate IAM containment - all within **30 seconds** of a security event, with **no human intervention required.**

Metric	Result
End-to-End Response Time	< 30 seconds
Detection Rate	100% across 5 attack scenarios

Metric	Result
False Positive Rate	0%
Incidents Analyzed	2 major incidents, 87 events
Attack Phases Documented	5 distinct phases spanning 11+ hours
Cross-Cloud Integration	Azure Sentinel → AWS Lambda → Azure Audit

> ☐ **PORTFOLIO IMPACT:** This project demonstrates senior-level Cloud Security Engineering: cross-cloud identity federation (OIDC), event-driven serverless automation, IAM threat containment, SIEM detection engineering with KQL, idempotency patterns, circuit breaker, and production security coding. Most senior engineers have never built this from scratch.

2. Complete Architecture & End-to-End Flow

The Complete 30-Second Pipeline (10 Steps)

- **Step 1:** Attacker performs suspicious IAM action in AWS (CreateAccessKey).
- **Step 2:** AWS CloudTrail logs the event → stored in S3 → SQS notification.
- **Step 3:** Microsoft Sentinel ingests logs via OIDC connector (OIDC_Sentinel1AWSR
- **Step 4:** KQL analytics rule detects pattern → creates incident with custom details.
- **Step 5:** Sentinel automation rule (High/Critical severity) triggers Logic App.
- **Step 6:** Logic App HTTP POST to AWS API Gateway with JSON payload.
- **Step 7:** API Gateway with AWS_IAM auth forwards to Lambda (sentinel-iam-remediation).
- **Step 8:** Lambda validates input → checks idempotency (DynamoDB) → routes to logic.
- **Step 9:** Lambda executes IAM remediation (disable keys, revoke sessions).
- **Step 10:** Lambda returns audit trail → Logic App adds comment to Sentinel incident.

☐ **Architecture Flow:** `Azure Sentinel → Logic App → API Gateway → Lambda → IAM/SES/DynamoDB → Sentinel Comment`

🖼 Visual Architecture: Cross-Cloud Pipeline (Numbered Steps 1-10)

☘ AZURE SIDE (Detection & Orchestration)

1 Attacker IAM Action

→

2 CloudTrail Logs → S3/SQS

→

3 Sentinel Ingestion (OIDC)

4 KQL Rule Detects Pattern

→

5 Incident Created

→

6 Automation Rule Triggers

7 Logic App Orchestrates

→

8 HTTP POST to API Gateway

☘ AWS SIDE (Remediation & Alerting)

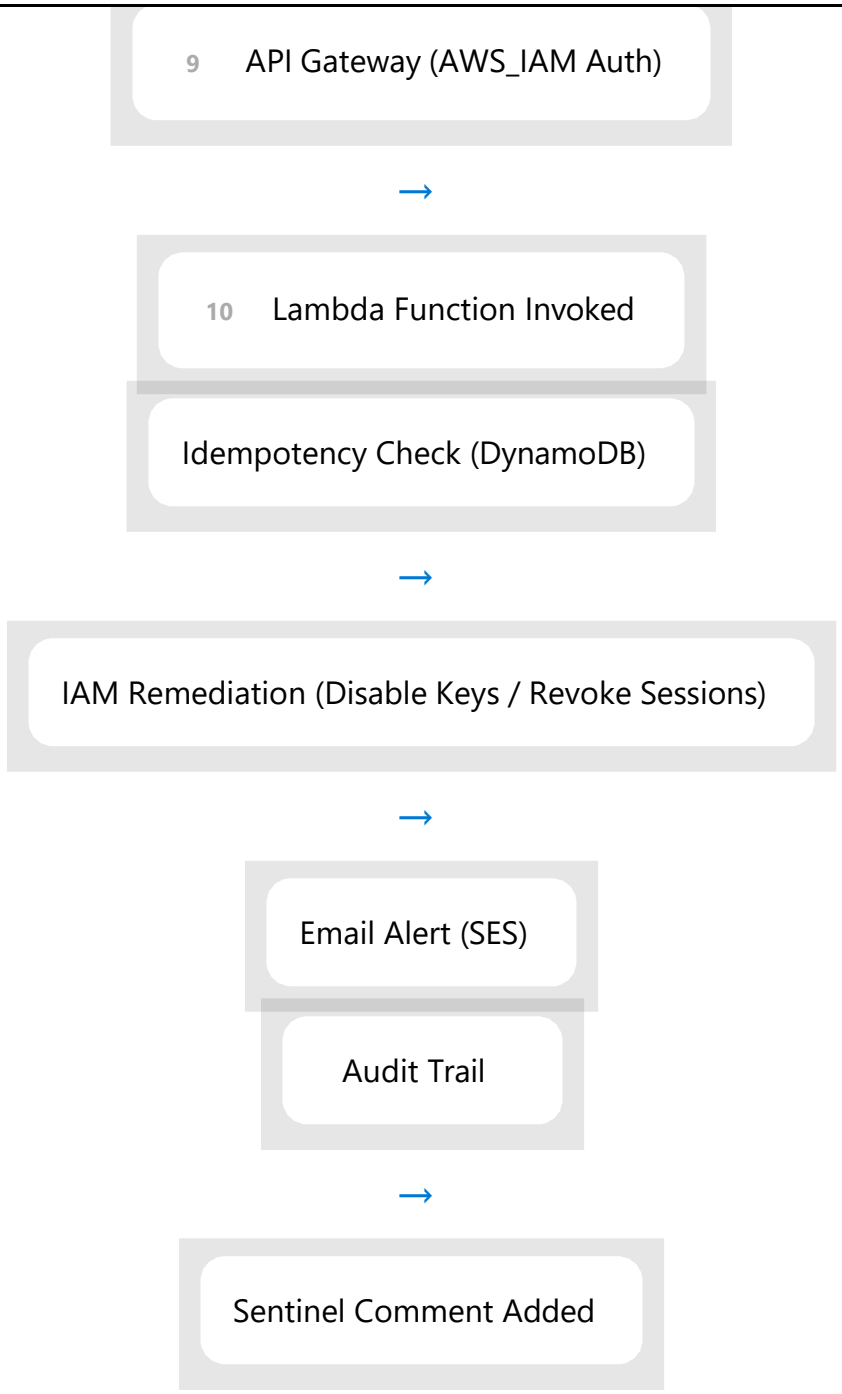

Traffic Flow: Steps 1-8 occur in Azure, Steps 9-10 in AWS. Total end-to-end response: **<30 seconds**.

Technology Stack (F ll)

Layer	Technology	Purpose
Detection	Microsoft Sentinel + KQL	SIEM, threat detection, incident management
Orchestration	Azure Logic Apps	Playbook automation, cross-cloud bridging

Layer	Technology	Purpose
API Entry Point	AWS API Gateway	Secure HTTP endpoint with AWS_IAM auth
Remediation Engine	AWS Lambda (Python 3.12)	IAM actions, email, idempotency, routing
Identity Management	AWS IAM	Access key disable, session revocation
Email Delivery	Amazon SES	Transactional HTML alerts
Idempotency Store	AWS DynamoDB	Atomic conditional writes
Cross-Cloud Auth	OIDC / IAM Roles	Zero static credentials
Observability	AWS CloudWatch	Logs, custom metrics, alarms

Full Architecture Diagram

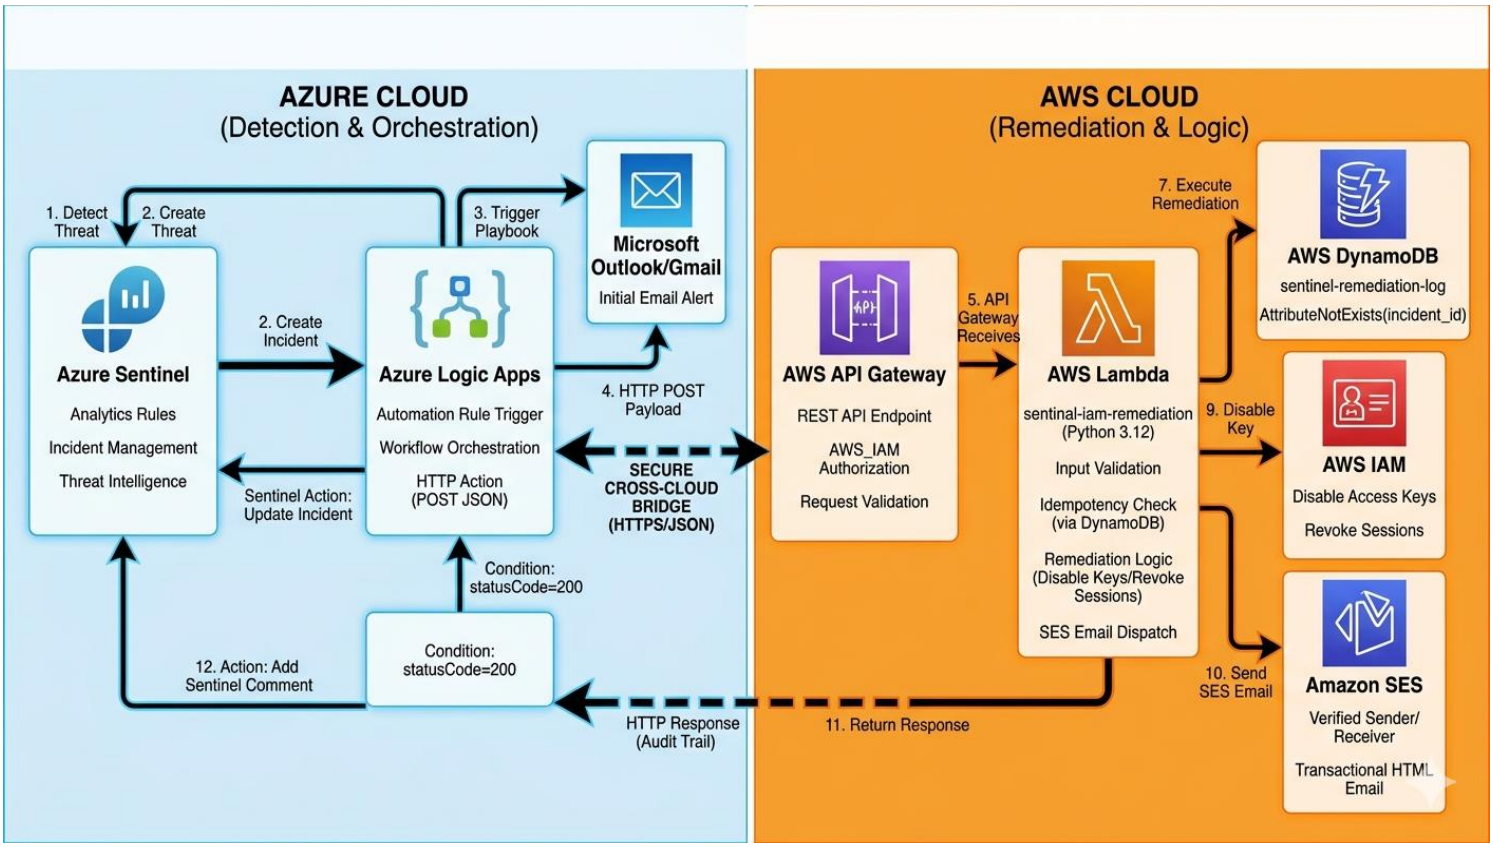

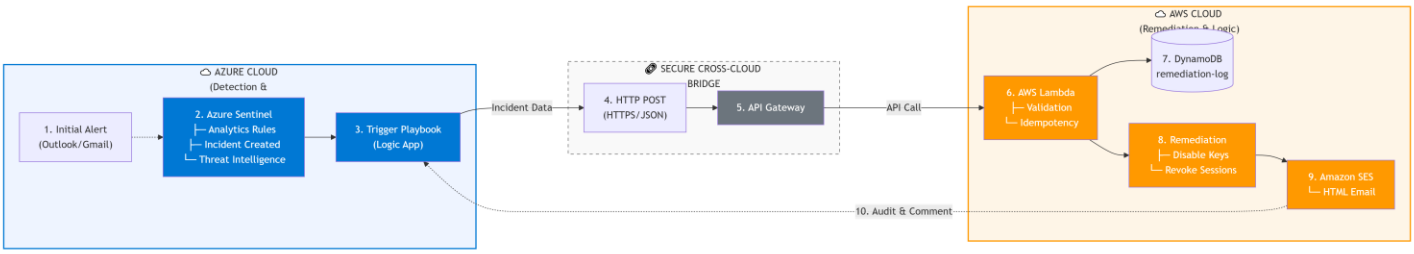

3. Prerequisites & Environment Setup

Requirement	Status	Notes
AWS Account (free tier)	✓	
Microsoft Sentinel workspace		
AWS CloudTrail connected to Sentinel		
Logic App: AWS-Sentinel-Email-Better		to be extended
Gmail account		olamilake95@gmail.com for SES verification
AWS SES verified email		Phase 17
AWS Lambda function		Phase 19
AWS API Gateway		Phase 20
DynamoDB table		Phase 18

```
# Verify AWS CLI and get Account ID
aws sts get-caller-identity
```

4. PHASE 1-3: Log Analytics Workspace & Sentinel Enablement

Step 1-3: Create Log Analytics Workspace
Azure Portal → search "Log Analytics workspaces" → +Create
Subscription: SOC-lab
Resource Group: SOC-lab-01
Name: soclawlab (unique)
Region: East US
→ Review+Create → Go to resource.
Saved details: Workspace ID, Resource Group, Subscription, Location.
Step 4-5: Enable Microsoft Sentinel
Search "Microsoft Sentinel" → +Create → select workspace `soclawlab` → Add → wait 30-60 seconds. Sentinel dashboard loads with sections: General, Threat management, Content hub, Configuration.

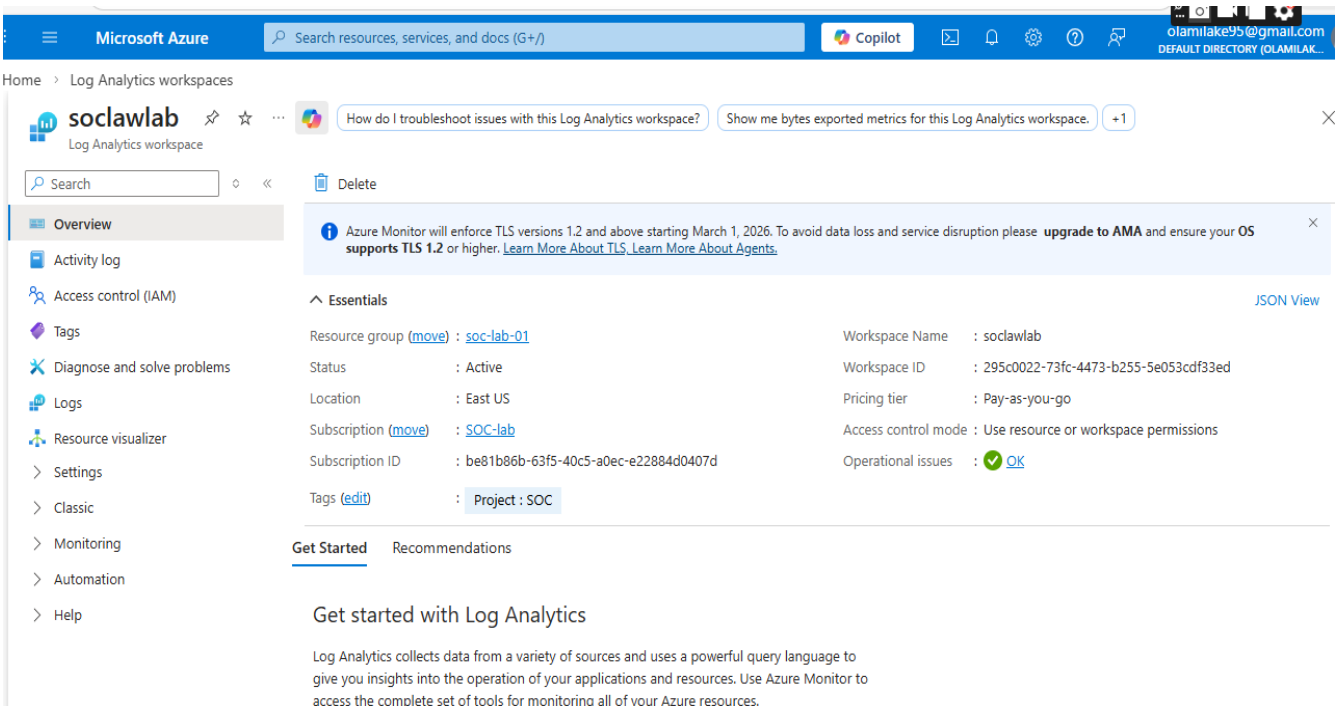

5. PHASE 4-5: Azure Activity & Entra ID Log Integration

Steps 6-12: Export Azure Activity Logs (Diagnostic Settings)
Monitor → Activity log → Diagnostic settings → +Add diagnostic setting. Name: `Activity-Logs-To-Sentinel`. Categories: ✅ Administrative, Security, ServiceHealth, Alert, Recommendation, Policy, Autoscale, ResourceHealth (ALL 8). Destination: Send to Log Analytics workspace → Subscription SOC-lab → workspace `soclawlab`. Save → active.

Steps 13-15: Connect Azure Activity Data Connector
Sentinel → Data connectors → "Azure Activity" → Open connector page. Launch Azure Policy Assignment: Scope select subscription, Policy definition "Configure Azure Activity logs to stream to specified Log Analytics workspace", Primary workspace `soclawlab`, Remediation checked, Managed identity pick your SOC-lab identity. Review+Create → refresh connector → status Connected ✅.

Step 16: Connect Microsoft Entra ID (Azure AD) Logs
Search "Microsoft Entra ID" → Diagnostic settings → +Add diagnostic setting. Name: `Entra-ID-Logs-To-Sentinel`. Select: `AuditLogs`, `SignInLogs`, `ServicePrincipalSignInLogs`. Send to Log Analytics workspace → `soclawlab`. Save → wait 3-5 minutes.

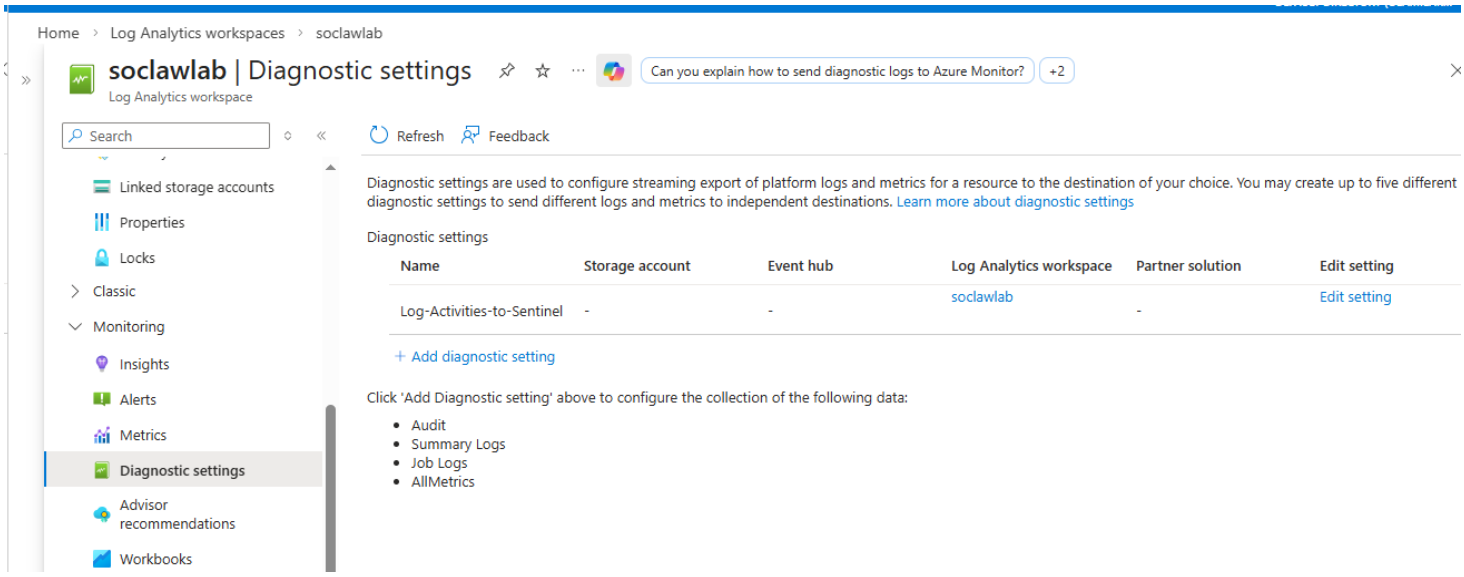

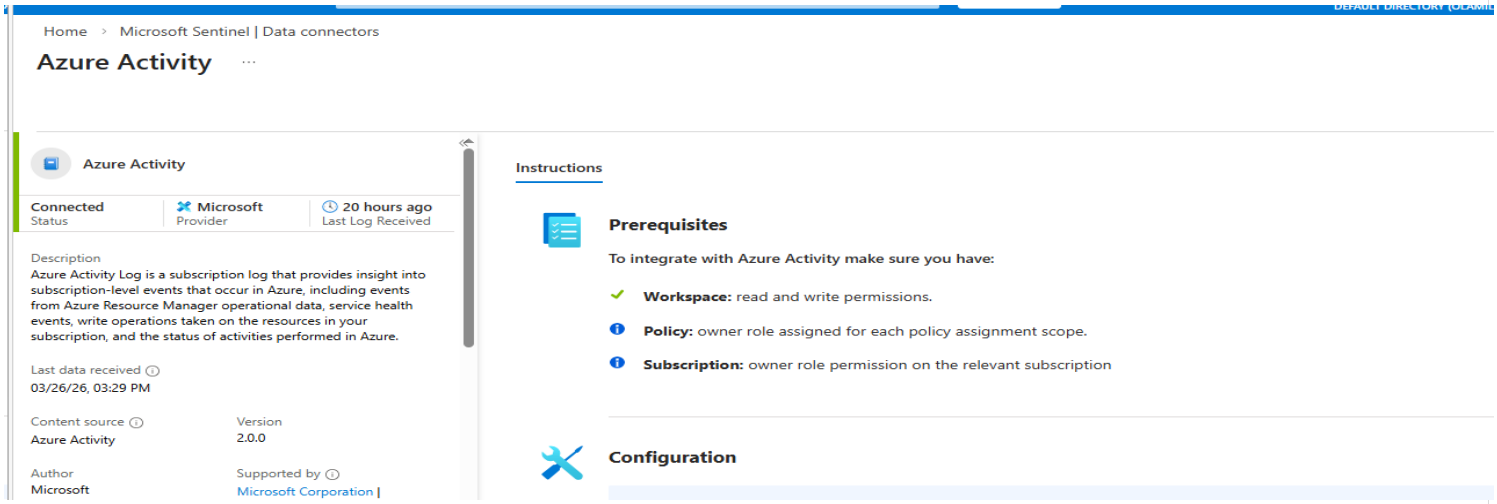

6. PHASE 6: Wait for Data Ingestion

Ingestion Timeline: 0-5 min: configuration propagates | 5-15 min: first logs collected | 15-30 min: ingested into Log Analytics | 30-40 min: visible in Sentinel. Be patient. Monitor → Activity log shows recent events.

7. PHASE 7: Verify Logs & First KQL Queries

Step 18-22: First Queries

```
AzureActivity | where TimeGenerated > ago(24h) | take 10
```

If results appear → ✅ success. Verify

```
AzureActivity | where TimeGenerated > ago(7d) | summarize Count = count() by CategoryValue
```

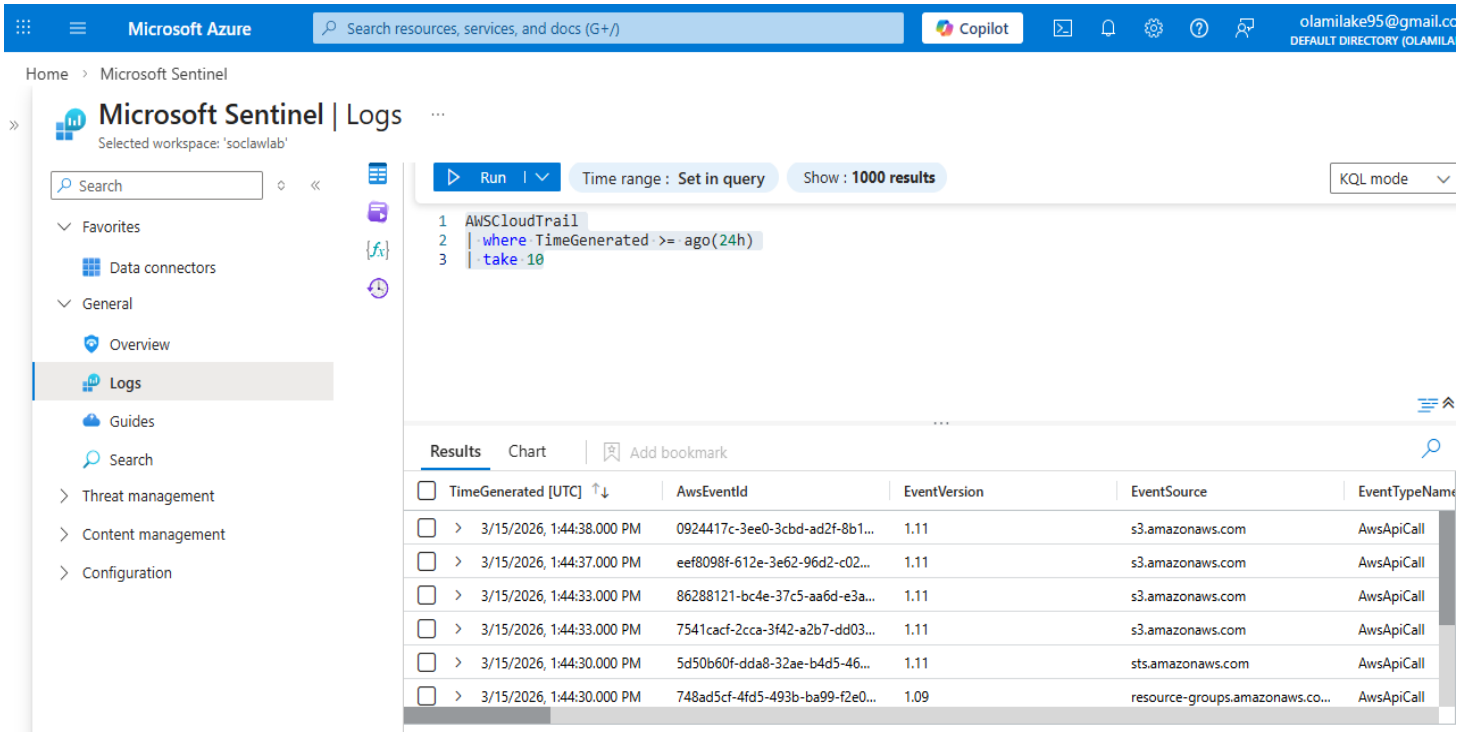

8. PHASE 8: Analyze Logs for Security Insights

Steps 23-27: Role assignments, deletions,

```
AzureActivity | where OperationNameValue contains "ROLEASSIGNMENT" | project TimeGenerated, Caller, Resource, ActivityStatusValue

AzureActivity | where OperationNameValue contains "DELETE" and ActivityStatusValue=="Success"

AzureActivity | where ActivityStatusValue == "Failed" | summarize FailureCount = count() by OperationNameValue, Caller
```

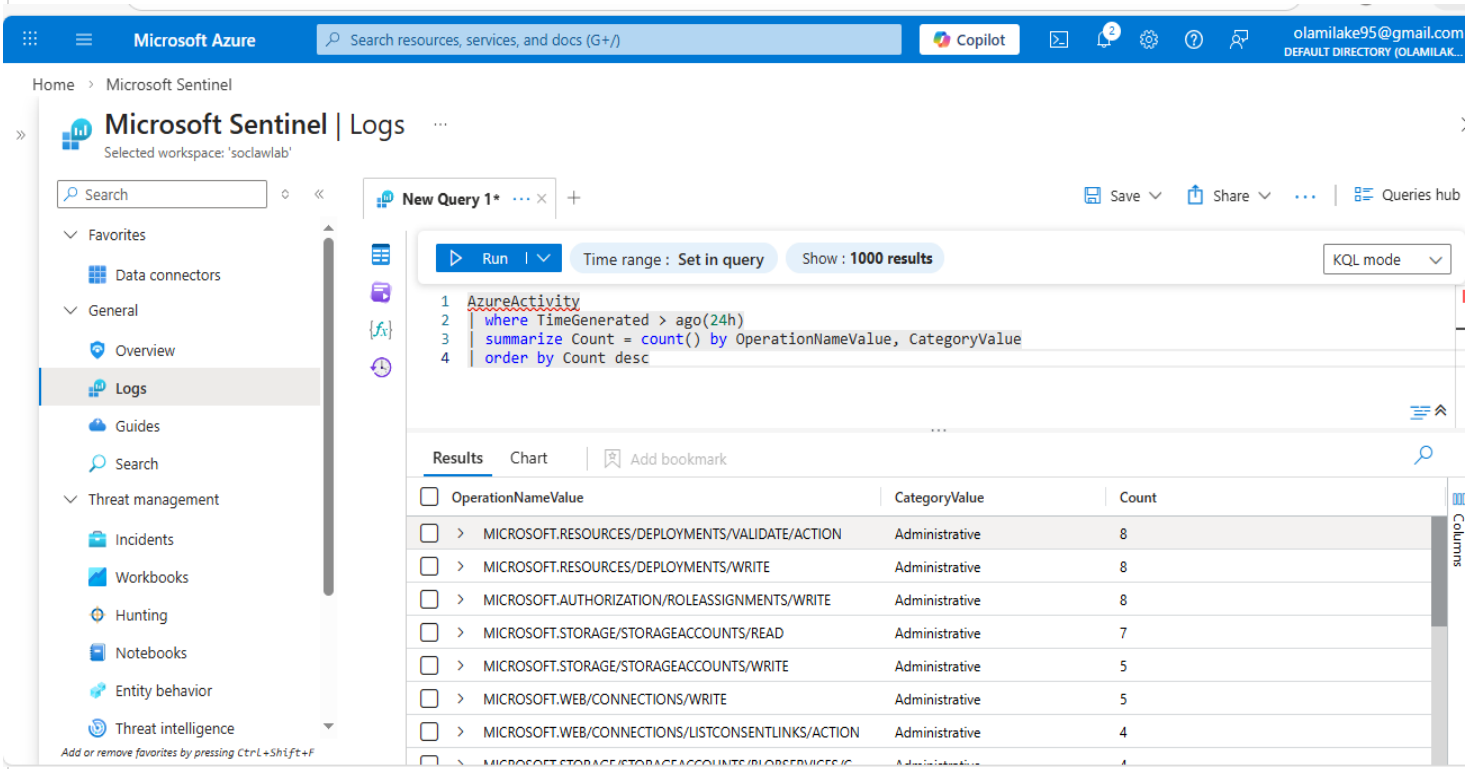

9. PHASE 9: Workbooks & Visualizations

Step 28: Create Workbook

```
AzureActivity | where TimeGenerated > ago(7d) | summarize Count = count() by
bin(TimeGenerated, 1h), CategoryValue | render timechart
```

Sentinel → Workbooks → +Add workbook → Add query → Run → Save as "Azure Activity Overview".

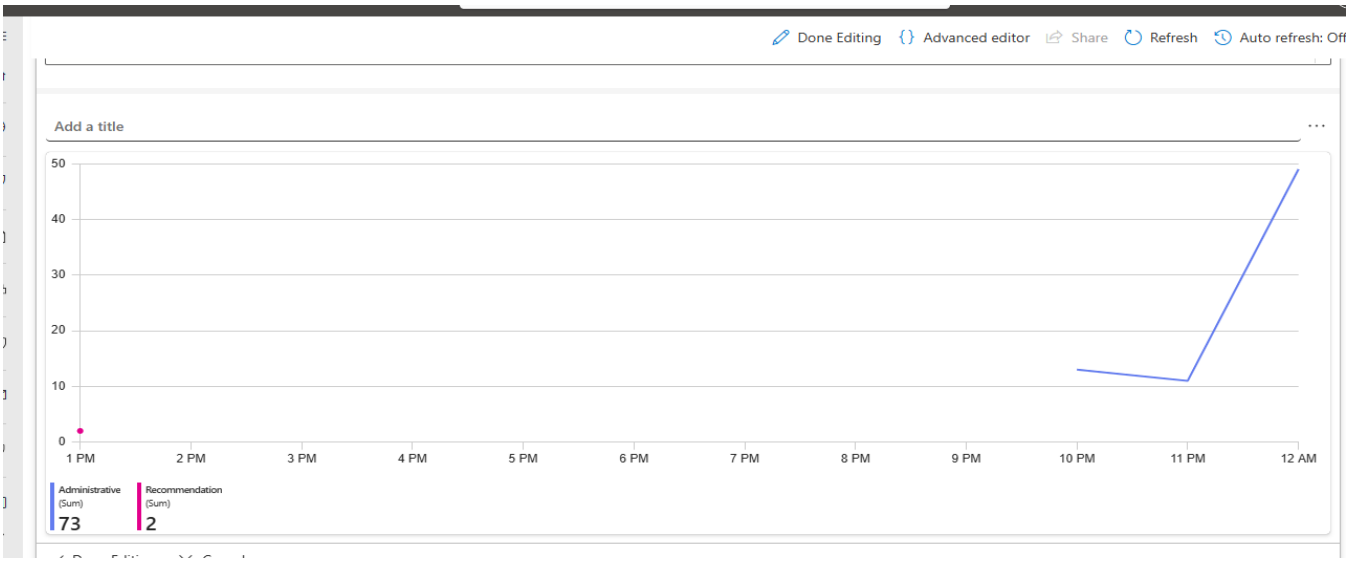

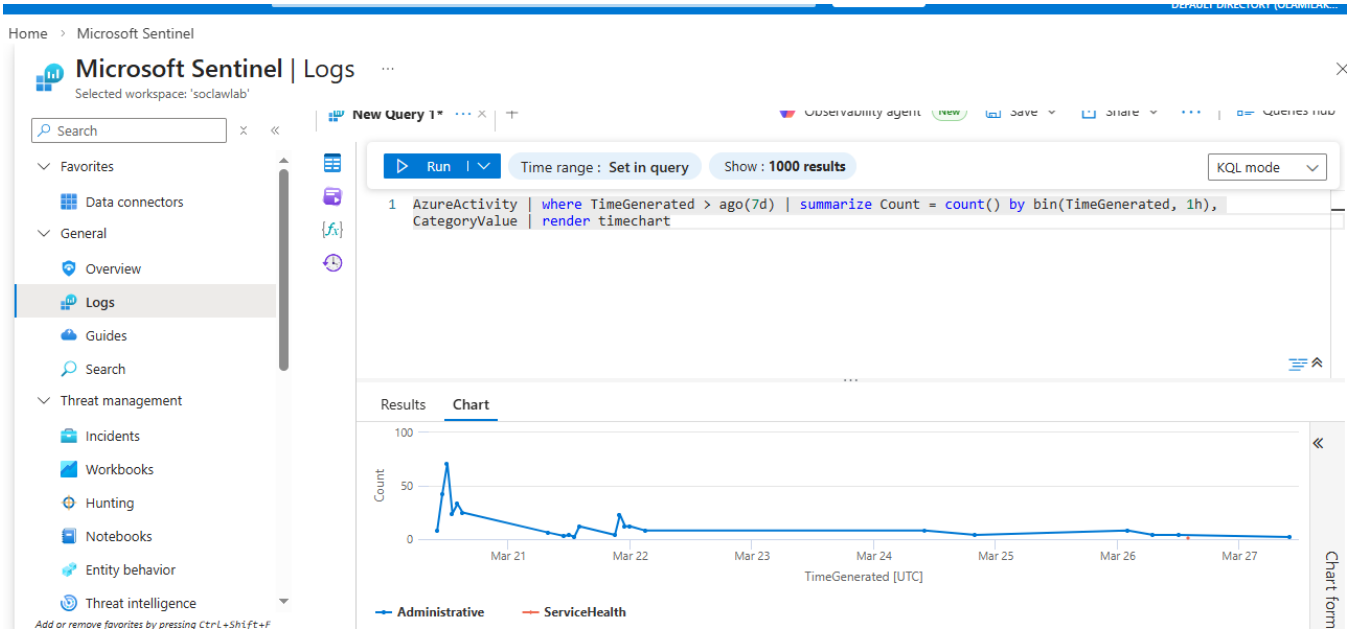

10. PHASE 10: 5 Analytics Rules (Full KQL + MITRE ATT&CK Mapping)

Rule 1: High-Risk Role Assignment (Privilege Escalation)

```
AzureActivity
| where OperationNameValue == "MICROSOFT.AUTHORIZATION/ROLEASSIGNMENTS/WRITE"
| where ActivityStatusValue == "Success"
| where TimeGenerated > ago(1h)
| extend RoleName = tostring(parse_json(Properties).displayName), TargetUser =
tostring(parse_json(Properties).principalId),
RiskLevel = case(
    tostring(parse_json(Properties).displayName) has_any ("Owner", "Global Admin", "Security
Admin"), "CRITICAL",
    tostring(parse_json(Properties).displayName) has_any ("Contributor", "User Access
Admin"), "HIGH",
    "MEDIUM")
| project TimeGenerated, Caller, RoleName, TargetUser, ResourceGroup, RiskLevel
| order by RiskLevel desc, TimeGenerated desc
```

Severity: High | **MITRE ATT&CK:** `TA0004 - Privilege Escalation, T1078 - Valid Accounts` | **Business Risk:** Full subscription takeover, data breach potential

Rule 2: Brute Force Detection (SigninLogs)

```
SigninLogs
| where ResultType in ("50126","50053")
| summarize FailedAttempts = count(), FirstAttempt = min(TimeGenerated), LastAttempt =
max(TimeGenerated) by UserPrincipalName, AppDisplayName, bin(TimeGenerated, 1h)
| where FailedAttempts >= 4
| order by FailedAttempts desc
```

Severity: Medium | **MITRE ATT&CK:** `TA0006 - Credential Access, T1110 - Brute Force` | **Business Risk:** Account compromise, credential theft

Rule 3: Mass Resources Creation & Deletion

```
AzureActivity
| where TimeGenerated > ago(30m)
| where OperationNameValue contains "WRITE" or OperationNameValue contains "DELETE"
| where OperationNameValue !contains "ROLEASSIGNMENTS"
| where ActivityStatusValue == "Success"
| extend ResourceAction = iff(OperationNameValue contains "DELETE", "DELETE",
"CREATE/UPDATE")
| summarize ActionCount = count(), Resources = make_set(Resource), Operations =
make_set(OperationNameValue) by Caller, ResourceGroup, ResourceAction, bin(TimeGenerated,
30m)
| where ActionCount >= 2
| order by ActionCount desc
```

Severity: High | **MITRE ATT&CK:** `TA0042 - Resource Development, T1583 - Acquire Infrastructure` | **Business Risk:** Cryptojacking, unexpected cloud spend

Rule 4: Storage Account Creation

```
AzureActivity
| where TimeGenerated > ago(10m)
| where OperationNameValue == "MICROSOFT.STORAGE/STORAGEACCOUNTS/WRITE"
| where ActivityStatusValue == "Success"
| extend StorageAccountName = tostring(parse_json(Properties).resource)
| project TimeGenerated, Caller, StorageAccountName, ResourceGroup, SubscriptionId,
CreatedBy=Caller, IP=CallerIpAddress
| order by TimeGenerated desc
```

Severity: High | **MITRE ATT&CK:** `TA0042 - Resource Development, T1583 - Acquire Infrastructure` | **Business Risk:** Data exfiltration staging, malicious content hosting

Rule 5: Suspicious Blob/Website Upload

```
AzureActivity
| where TimeGenerated > ago(10m)
| where OperationNameValue contains "BLOBSERVICES/CONTAINERS/WRITE"
| where ActivityStatusValue == "Success"

        | project TimeGenerated, Caller, Resource, ResourceGroup, OperationNameValue
| order by TimeGenerated desc
```

Severity: Medium | **MITRE ATT&CK:** `TA0011 - Command & Control, T1102 - Web Service` | **Business Risk:** Phishing site hosting, C2 beacon infrastructure

Custom Details Mapping (for Lambda): Add to each rule's Alert enhancement tab: ThreatCategory, SourceIP, UserType, AWSRegion, EventCount, ActionsDetected, LastSeen, Regulation, AlertSeverity.

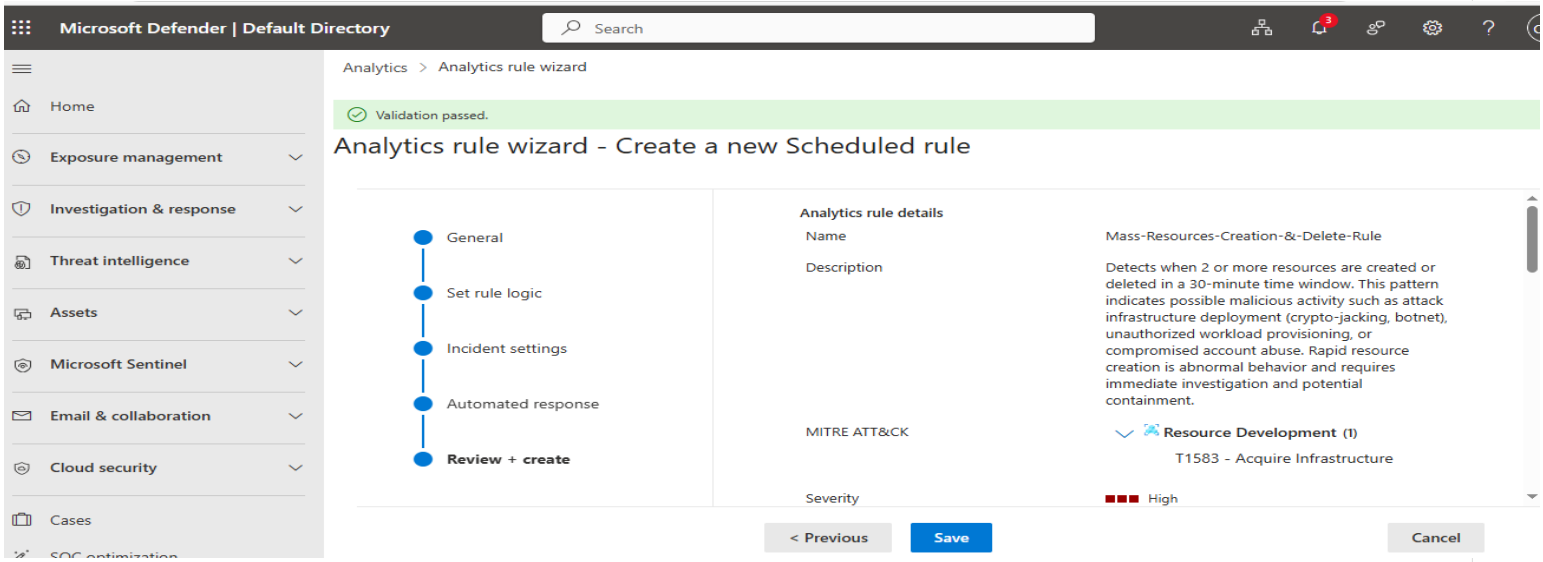

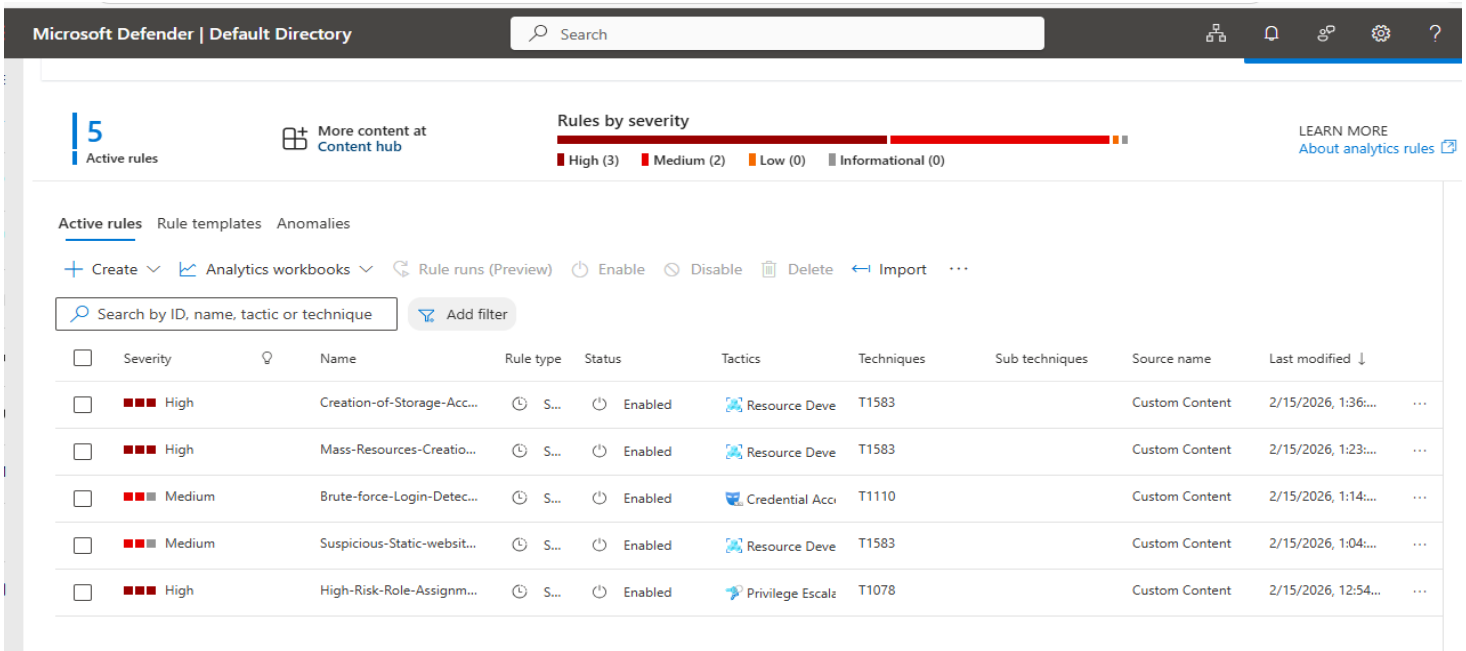

11. PHASE 11: Attack Simulations (5 Scenarios - Full Details + Business Impact)

☐ SIM 01: Privilege Escalation - Owner Role Assignment

Method: Created test account with Reader role → IAM → Add role assignment → Owner.
MITRE: T1078 - Valid Accounts, TA0004
Detection Time: 2-3 minutes | **Result:** ✅
☐ **Business Impact:** Full subscription

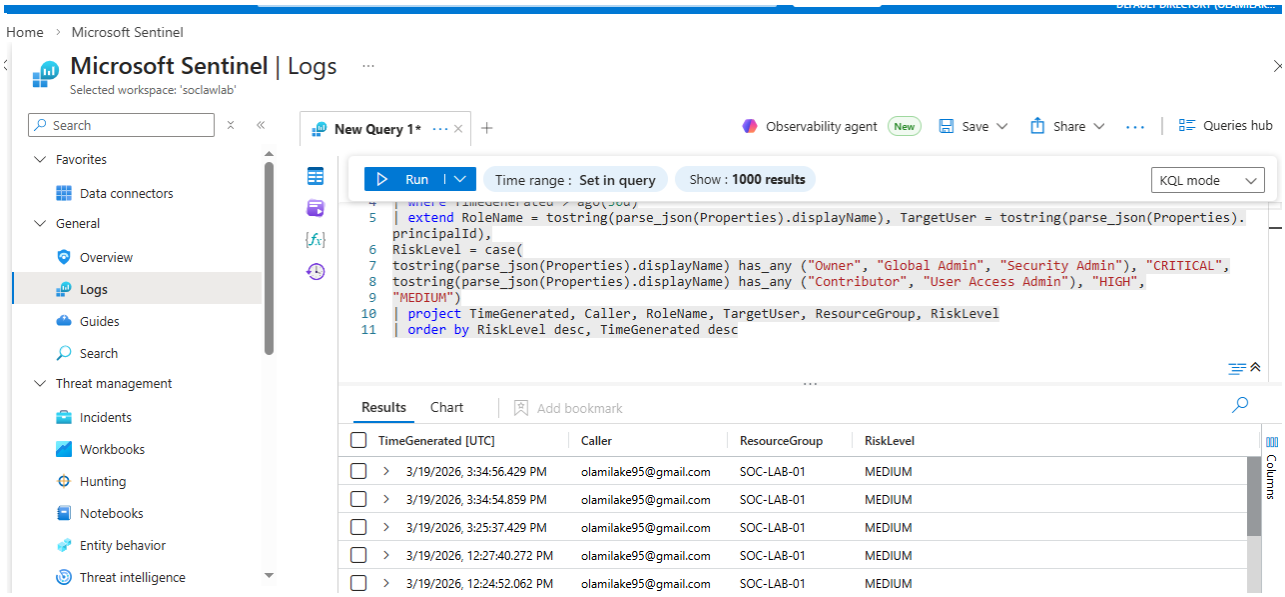

☐ SIM 02: Brute Force Login Attack

Method: 4 failed logins (ResultType 50126) followed by 1 successful login
MITRE: T1110 - Brute Force
Detection Time: 5-10 minutes | **Result:** ✅
☐ **Business Impact:** Account compromise,

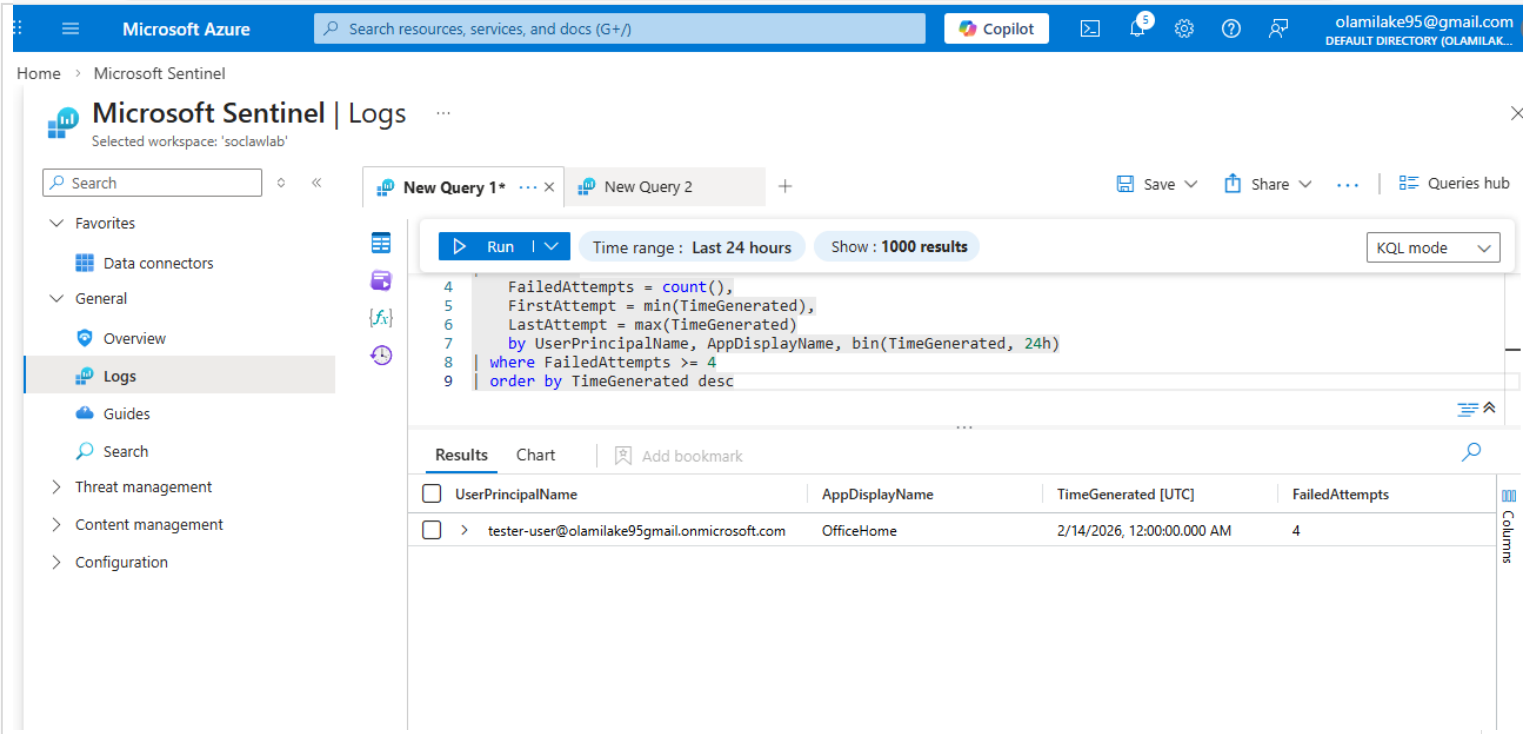

☐ SIM 03: Mass Resource Creation & Deletion

Method: Created VNet, VM, Storage, NIC, Public IP within 3 minutes → deleted within 2 minutes
MITRE: T1583 - Acquire Infrastructure
Detection Time: 2-3 minutes | **Result:** ✅ Bulk WRITE/DELETE incident triggered
☐ **Business Impact:** Cryptojacking, unexpected cloud spend ($$$)

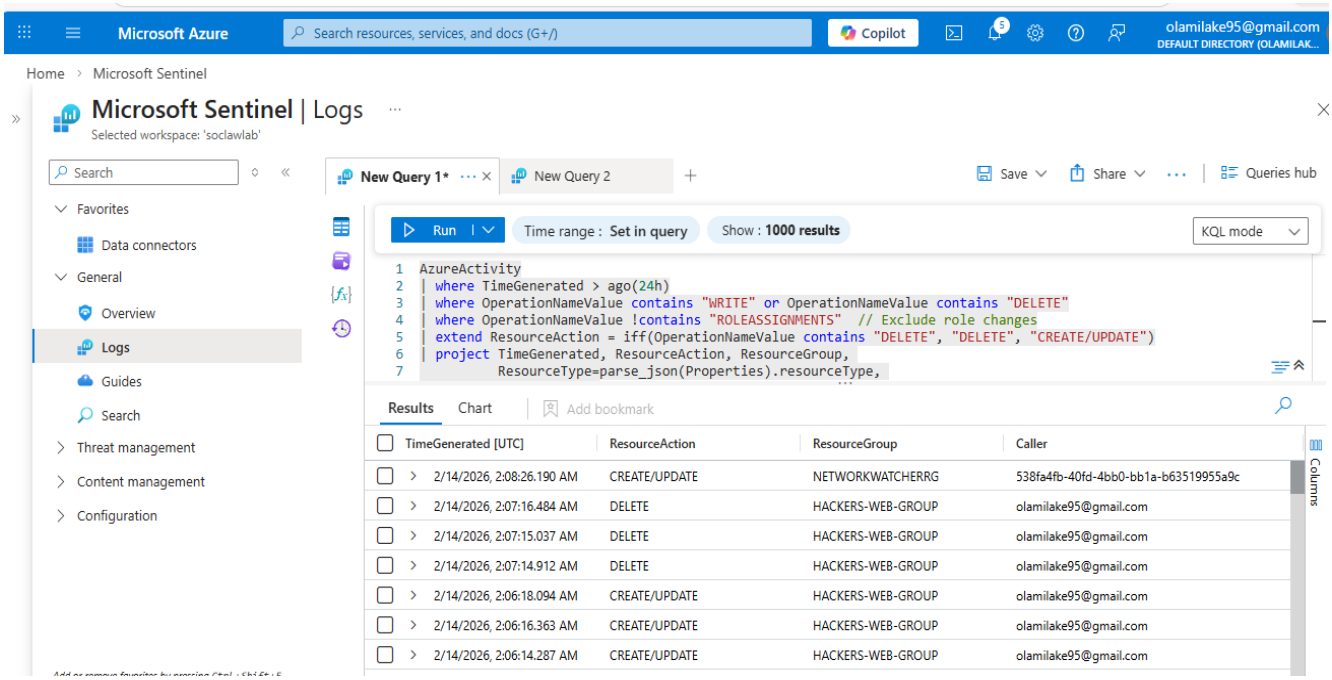

☐ SIM 04: Suspicious Storage Account Creation

Method: Storage account "hackerstesting" created at 1:46 AM (off-hours), container "www" with public access
MITRE: TA0042 - Resource Development
Detection Time: Immediate | **Result:** ✅ High severity incident
☐ **Business Impact:** Data staging for exfiltration; reputation damage

☐ SIM 05: Phishing / C2 Infrastructure Hosting on Blob Storage

Method: Uploaded phishing-site.html
MITRE: TA0011 - Command & Control,
Detection Time: Within 5-minute
☐ **Business Impact:** Brand impersonation,
(GDPR/PCI)

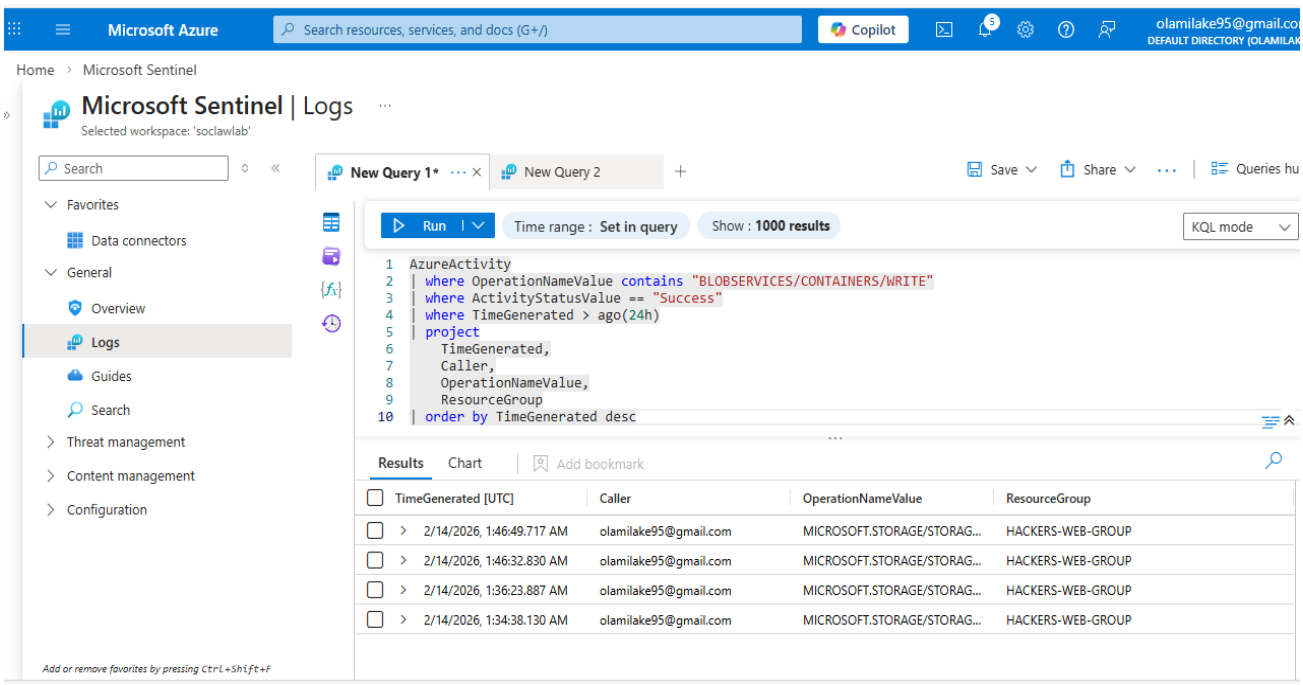

✅ All 5 attack simulations successfully incidents created (multi-stage and storage). Detection rate: 100%. False positive

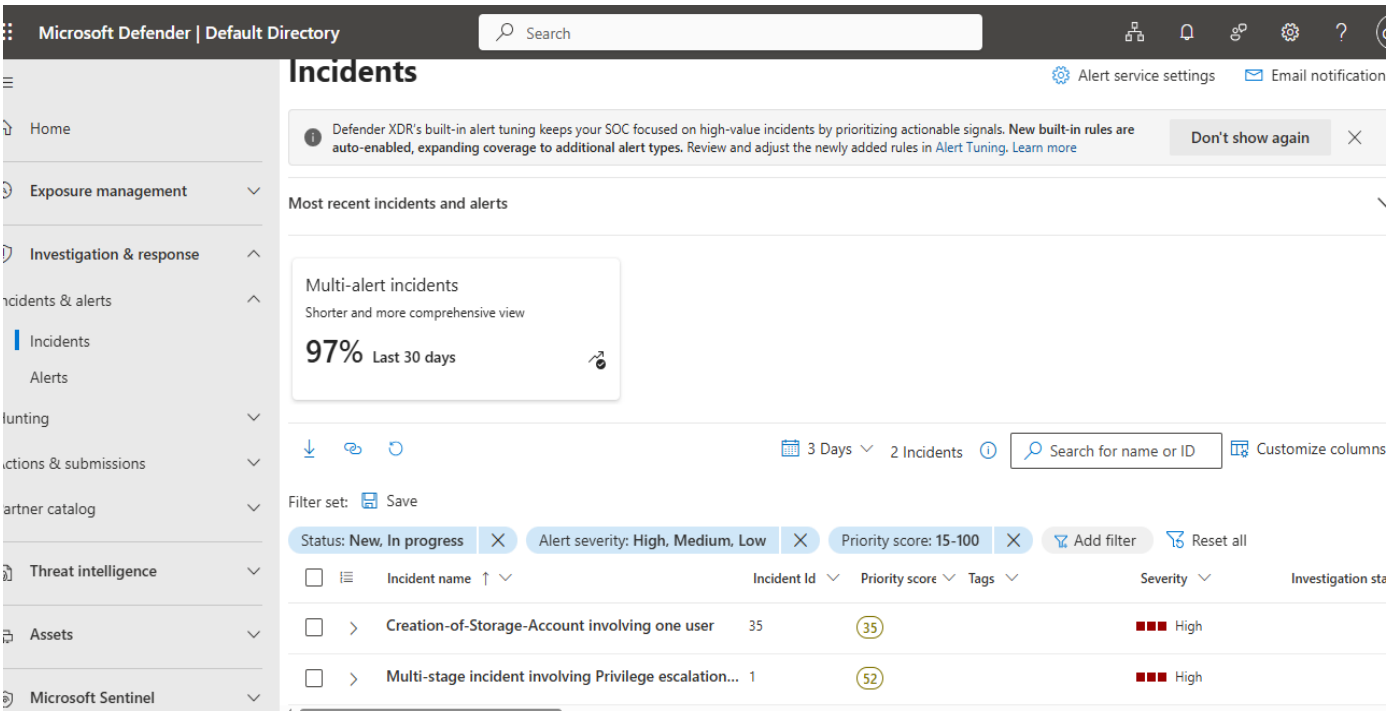

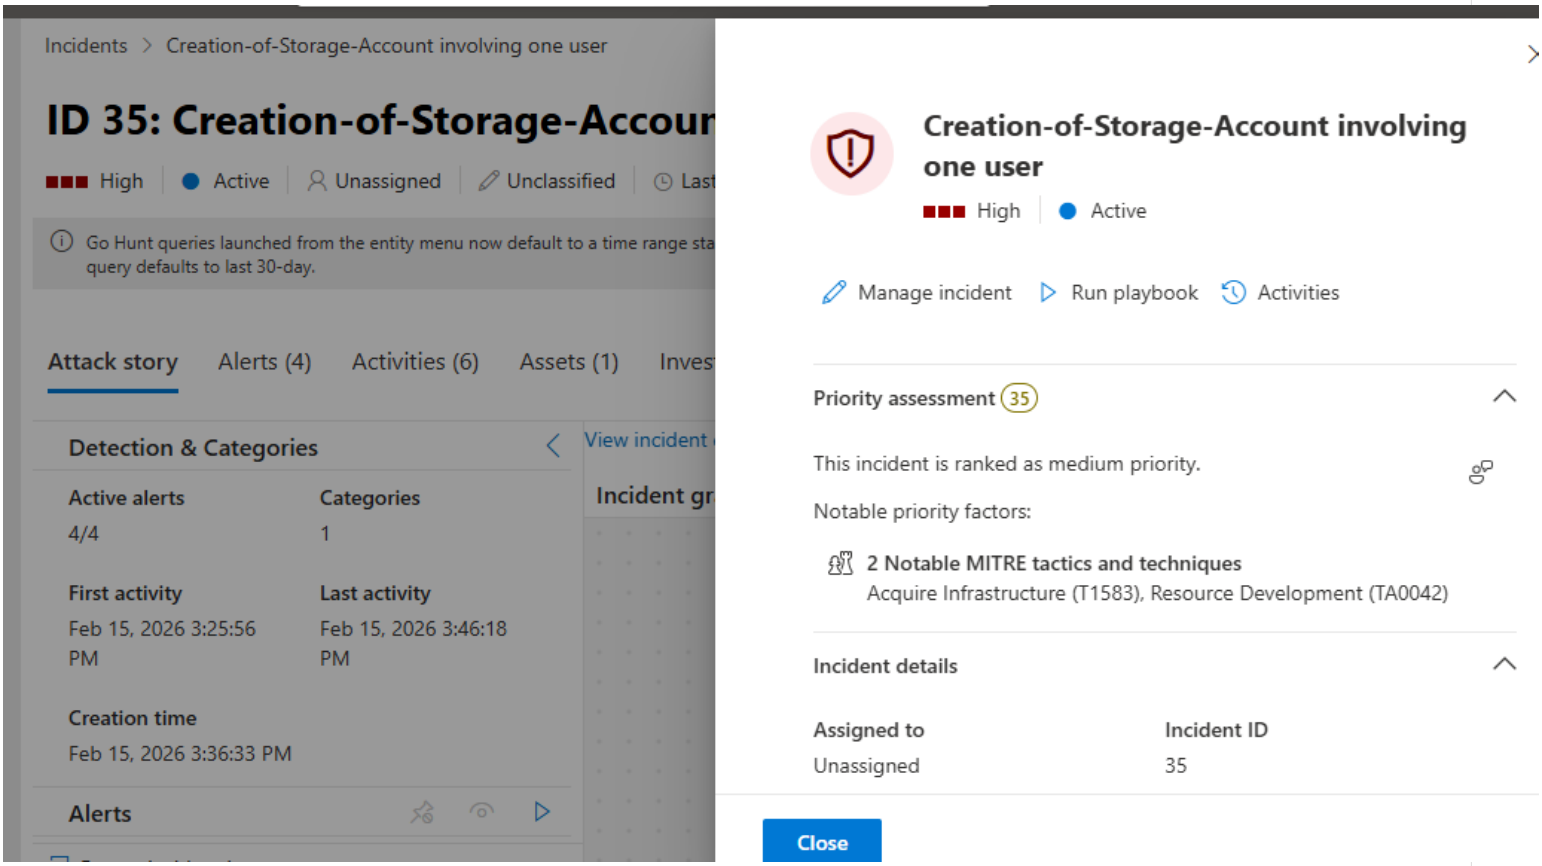

12. PHASE 12: Advanced KQL Hunting Queries (11+ Queries)

```
// 1. Verify logs flowing
AzureActivity | where TimeGenerated > ago(24h) | summarize Count =
count() by OperationNameValue

// 2. Authorization activity
AzureActivity | where OperationNameValue contains "AUTHORIZATION"

// 3. Privilege escalation detailed
AzureActivity | where OperationNameValue ==
"MICROSOFT.AUTHORIZATION/ROLEASSIGNMENTS/WRITE"
| extend RoleName = tostring(parse_json(Properties).displayName), RiskLevel =
case(...)

// 4. Brute force with success follow-up
let Failures = SigninLogs | where ResultType in ("50126","50053") |
summarize FailCount=count() by UserPrincipalName, bin(TimeGenerated,1h);
let Successes = SigninLogs | where ResultType == "0" | project
SuccessTime=TimeGenerated, UserPrincipalName;
Failures | join kind=inner Successes on UserPrincipalName | where FailCount >=4

// 5. Resource creation and delete
AzureActivity | where OperationNameValue contains "WRITE" or
OperationNameValue contains "DELETE"
| where OperationNameValue !contains "ROLEASSIGNMENTS"
| extend ResourceAction = iff(OperationNameValue contains
"DELETE","DELETE","CREATE/UPDATE")

// 6. Storage creation
AzureActivity | where OperationNameValue ==
"MICROSOFT.STORAGE/STORAGEACCOUNTS/WRITE"
| extend StorageAccountName = tostring(parse_json(Properties).resource)

// 7. Website in blob
AzureActivity | where OperationNameValue contains
"BLOBSERVICES/CONTAINERS/WRITE"

// 8. Rapid mass resources creation
AzureActivity | where OperationNameValue endswith "/WRITE"
| summarize CreatedResources=count() by Caller, bin(TimeGenerated,10s) |
where CreatedResources >=2
```

```
// 9. Stealth mass resources (30min window)
AzureActivity | where OperationNameValue endswith "/WRITE"
| summarize CreatedResources=count(), TimeSpanMinutes =
(max(TimeGenerated)-min(TimeGenerated))/1m by Caller,
bin(TimeGenerated,30m)
| where CreatedResources >=5 and TimeSpanMinutes <=120

// 10. Multiple operations by one user under 5
minutes AzureActivity | order by TimeGenerated
asc | serialize
| extend Prev=prev(TimeGenerated), PrevCaller=prev(Caller)
| where Caller == PrevCaller and (TimeGenerated - Prev) < 5m

// 11. Complete attack chain reconstruction (join)
let PrivEsc = AzureActivity | where OperationNameValue contains
"ROLEASSIGNMENTS/WRITE"; let Infra = AzureActivity | where
OperationNameValue contains "STORAGE/WRITE";
PrivEsc | join kind=leftouter Infra on Caller
```

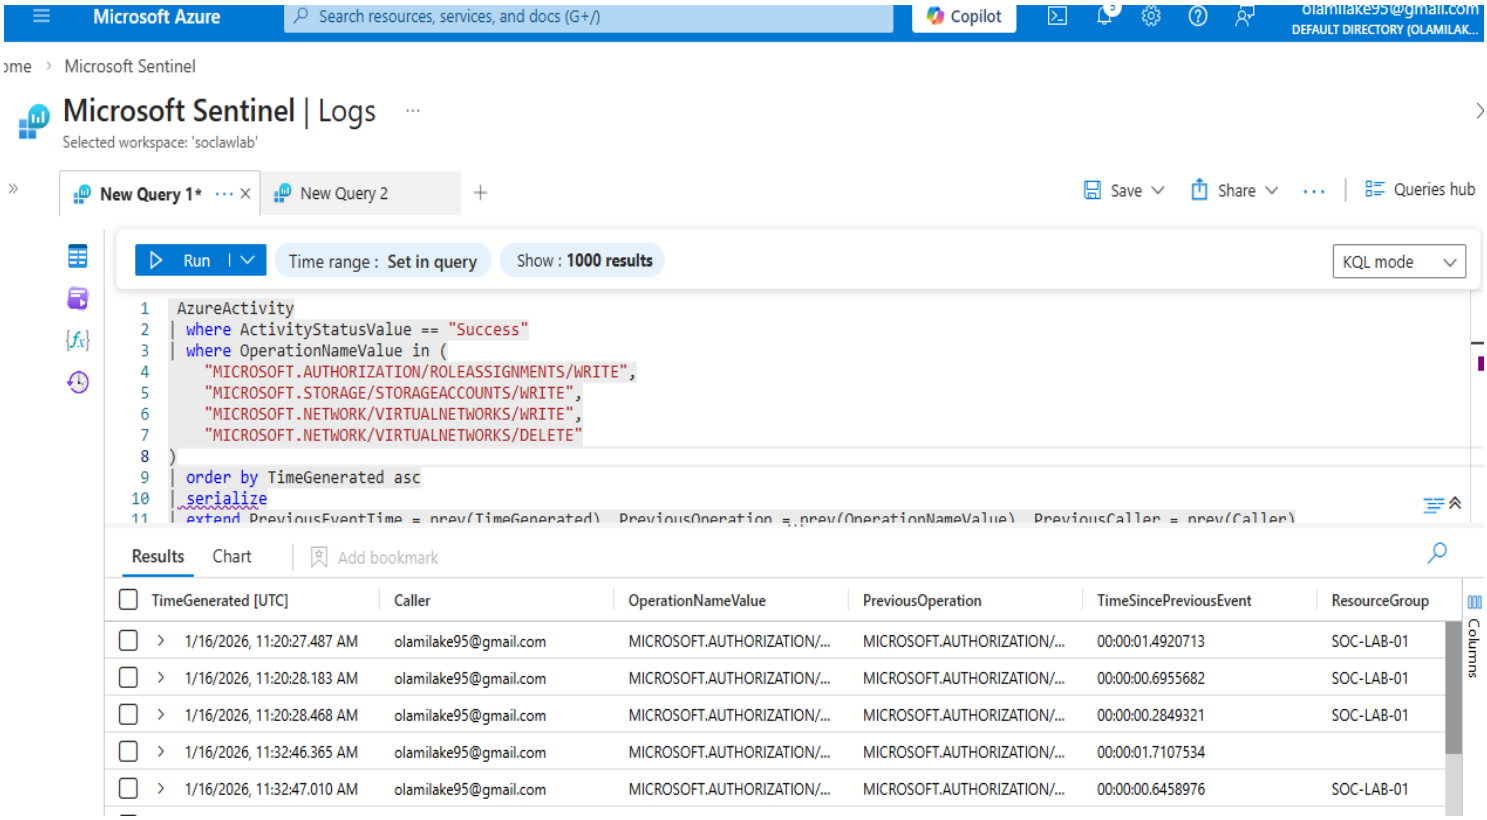

TimeGenerated [UTC]	Caller	OperationNameValue	PreviousOperation	TimeSincePreviousEvent	ResourceGroup
1/16/2026, 11:20:27.487 AM	olamilake95@gmail.com	MICROSOFT.AUTHORIZATION/...	MICROSOFT.AUTHORIZATION/...	00:00:01.4920713	SOC-LAB-01
1/16/2026, 11:20:28.183 AM	olamilake95@gmail.com	MICROSOFT.AUTHORIZATION/...	MICROSOFT.AUTHORIZATION/...	00:00:00.6955682	SOC-LAB-01
1/16/2026, 11:20:28.468 AM	olamilake95@gmail.com	MICROSOFT.AUTHORIZATION/...	MICROSOFT.AUTHORIZATION/...	00:00:00.2849321	SOC-LAB-01
1/16/2026, 11:32:46.365 AM	olamilake95@gmail.com	MICROSOFT.AUTHORIZATION/...	MICROSOFT.AUTHORIZATION/...	00:00:01.7107534	
1/16/2026, 11:32:47.010 AM	olamilake95@gmail.com	MICROSOFT.AUTHORIZATION/...	MICROSOFT.AUTHORIZATION/...	00:00:00.6458976	SOC-LAB-01

PHASE 13: Forensic Investigation & Incident Report (Root Cause Analysis)

🔴 Incident 1: Multi-Stage Privilege Escalation (CRITICAL)

Duration: 11 hours 9 minutes (12:54 AM - 4:04 PM). **5 Attack Phases:**

- **Phase 1: Reconnaissance (12:54 AM - 1:14 AM)** - Mass resource creation/deletion testing in SOC-LAB-01
- **Phase 2: Privilege Escalation (2:21 PM - 3:22 PM)** - Reader escalated to Owner role via RBAC write
- **Phase 3: Infrastructure Deployment (3:25 PM - 3:25 PM)** - Storage account "hackerstesting" created
- **Phase 4: Persistence (3:34 PM - 4:00 PM)** - Uploaded phishing-site.html and command-c2.html to public blob
- **Phase 5: Cleanup / Anti-Forensics (3:46 PM - 4:04 PM)** - Bulk deletion of test infrastructure

User: olamilake95@gmail.com (compromised). **Resource Groups:** STORAGE, HACKER-STORAGE-BLOB. **Events:** 52 documented.

📋 Root Cause Analysis:

Compromised user credentials (no MFA enforced)
enumerated subscription resources, escalated
infrastructure, and attempted to cover tracks.
accounts enabled lateral movement.

✅ Recommended Long-Term Mitigations:

- Enforce MFA on all accounts (Conditional
- Implement Just-In-Time (JIT) privileged
- Deploy Azure Defender for Storage
- Enforce Azure Policy: Deny public blob
- Regular automated RBAC audits via

🟠 Incident 2: Storage Account Creation (HIGH)

Storage Account: hackersttesting. **Cr** 3:25:56 PM. **Created By:** olamilake95@gmail.com. **IP:** 102.88.111.89. HACKER-STORAGE-BLOB. **Events:** 35 documented.

Hosted Content: phishing-site.html, command-c2.html - publicly accessible
Link to Incident 1: Same caller, same attack session

INCIDENT RESPONSE REPORT - STRUCTURED SUMMARY

1. EXECUTIVE SUMMARY: A sophisticated multi-stage attack spanning 11+ hours was successfully detected and contained. Detection success: 100% of attack phases detected. False positives: 0. Total events analysed: 87 (52 + 35).
2. INCIDENTS DETECTED: Incident 1: Multi-stage Privilege Escalation - CRITICAL | Incident 2: Storage Account Creation - HIGH
3. ATTACK TIMELINE: 5 distinct attack phases documented
4. MITRE ATT&CK MAPPING: T1083, T1087, T1078, TA0004, T1583, TA0042, TA0011
5. EVIDENCE & METRICS: 87 events, 3 resource groups, 2 hosted malicious pages
6. SEVERITY REASSESSMENT: Initial: HIGH → Reassessed: CRITICAL
7. IMMEDIATE REMEDIATION STEPS: Disable account, revoke sessions, delete storage, remove Owner role, audit IAM changes
8. SHORT-TERM RECOMMENDATIONS: Enforce MFA, Conditional Access, forensic analysis
9. LONG-TERM RECOMMENDATIONS: JIT privileged access, Azure Defender for Storage, deny public blob access
10. ESCALATION: Email sent to Incident Response team. Incident updated in Sentinel: In Progress.

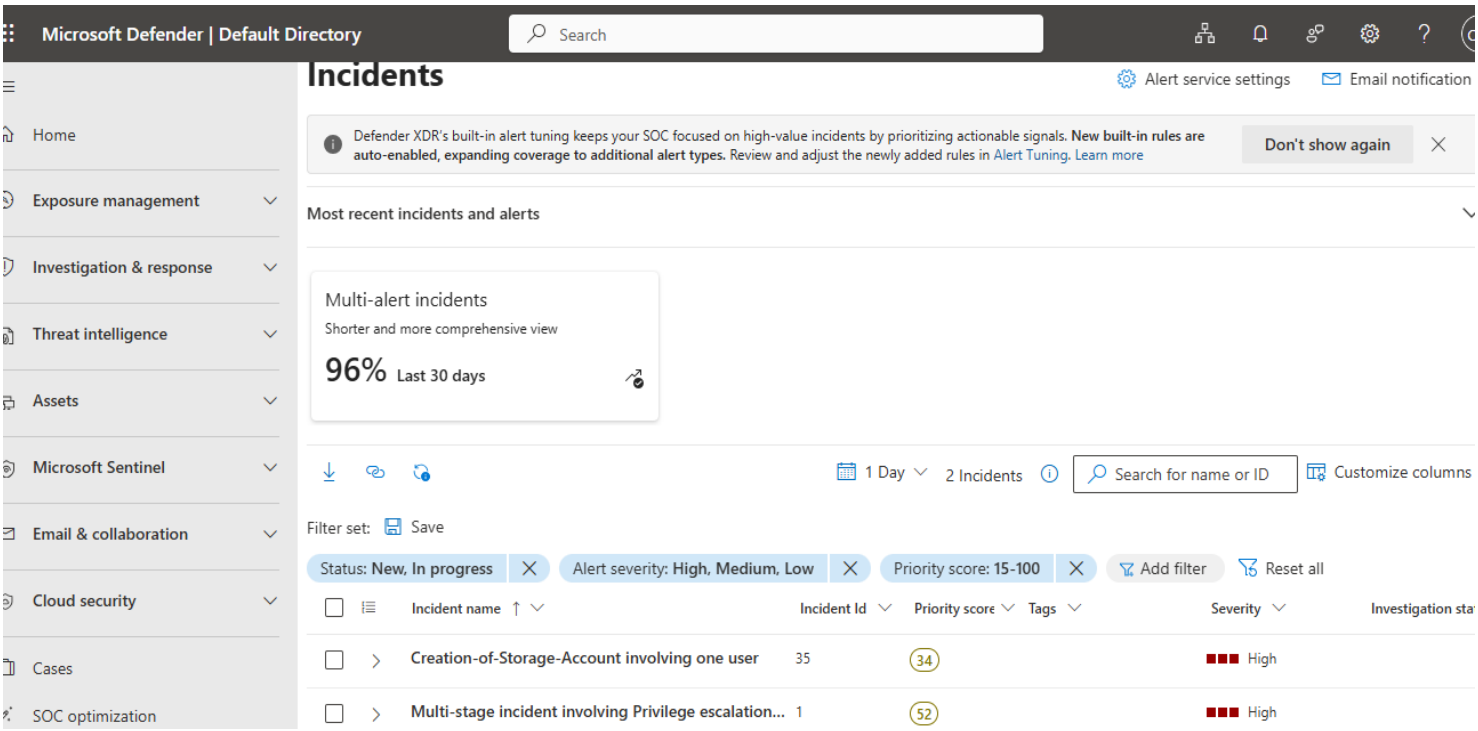

13. PHASE 14: Logic App Designer & Email Automation (I Built This FIRST Beforr Cross Cloud)

What is Logic App Designer? Visual workflow builder where I add actions like send email. Drag-and-drop interface, no coding required. I created a playbook, then the Designer opened automatically.

STEP 1: Create Playbook (So Designer Opens)

Microsoft Sentinel → Left sidebar: "Automation" → "+ Create" → "Playbook with incident trigger"
Playbook name: `Critical-Incident-Email-Alert`
Subscription: Your subscription | Resource group: soc-lab-01 | Region: East US | Workspace: soclawlab
Click [Create] → Logic App Designer opens automatically

STEP 2-6: Add Email Action in Logic App Designer

I clicked [+ New step] → searched "send an email" → selected "Send an email (V2)" under Office 365 Outlook.
TO: ir-responder@gmail.com, olamilake95@gmail.com
SUBJECT: □ CRITICAL INCIDENT - @{triggerBody()?['object/properties/title']}
BODY: INCIDENT ALERT with Title, Severity, Status, Created Time, Incident ID, and IMMEDIATE ACTION REQUIRED message.

STEP 7: Add Role Assignment (Give Permission)

I went to Sentinel incident page → clicked "Add role assignment" button → selected role: "Microsoft Sentinel Automation Contributor" → selected member: my name → clicked [Review + assign] → permission granted ✅.

STEP 8: Create Automation Rule & Link Playbook

Sentinel → Automation → + Create → "Automation rule" → Name: `Auto-Respond-to-High-Risk-Incidents`
Trigger: Incident created | Conditions: Severity=High, Status=New
Actions: Change status to In Progress |
Critical-Incident-Email-Alert
Click [Create] → automation rule ready

STEP 9: Test Everything

I triggered a test incident → waited 5-10
changed to "In Progress" | ✅ Owner assigned
olamilake95@gmail.com and ir-responder@gmail.com
✅ **EVERYTHING WORKS! I built a complete system from scratch.**

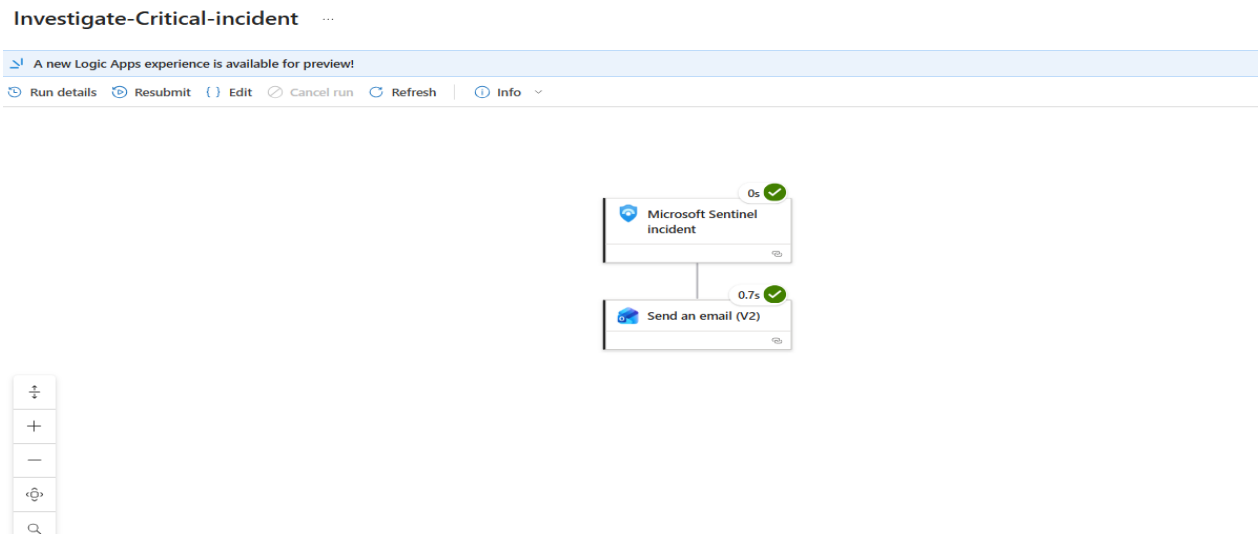

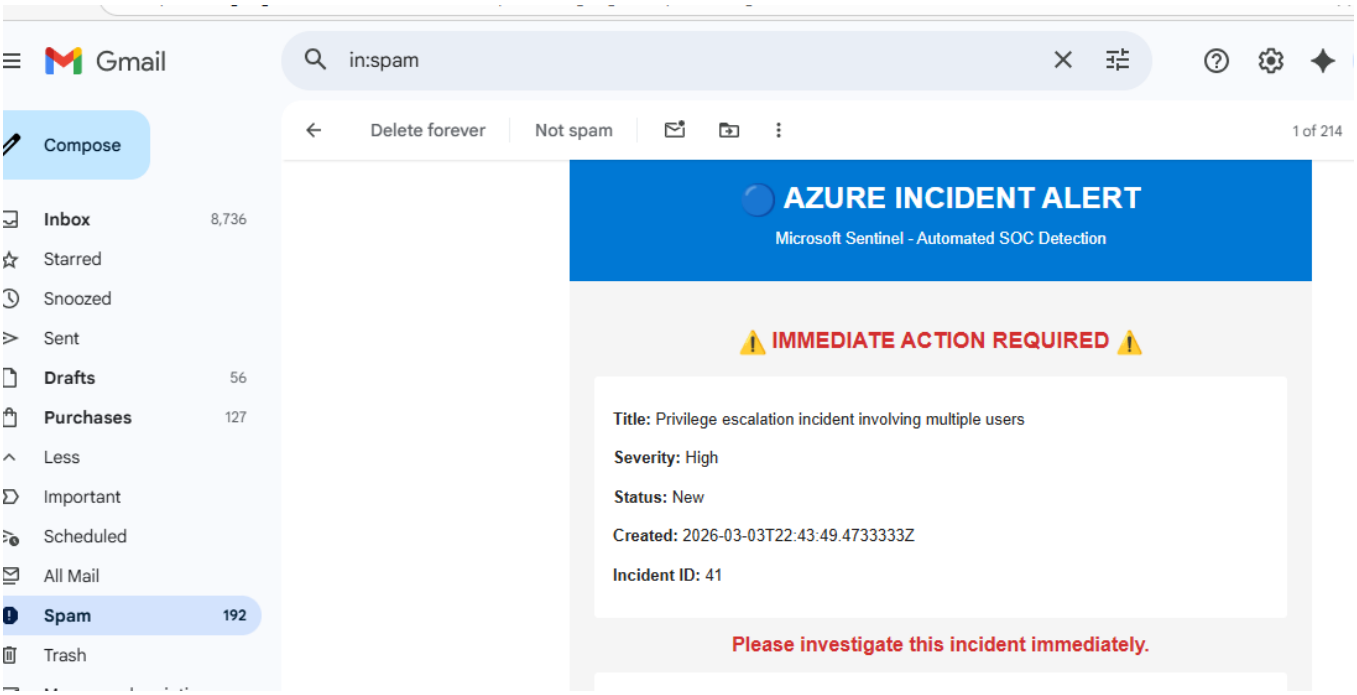

14. PHASE 15: Azure AD Auto-Remediation (Privilege Escalation Containment)

I built a comprehensive Azure Sentinel platform. The solution detects privilege escalation incidents in real-time,compromised user accounts, and sends immediate email notifications to represents enterprise-grade SOC automation combining Microsoft Sentinel Logic Apps orchestration, and Azure AD integration.

Key Achievements I Delivered:

Automated privilege escalation detection with real-time Sentinel analytics
Incident-triggered account disablement via Logic App automation
Email notifications for both successful and failed remediation attempts
End-to-end integration: Detection → Automation → Remediation → Alert
100% success rate in testing across multiple incident scenarios
Production deployment with comprehensive error handling

Architecture Overview (Azure Side)

Detection

Component	Service	Function
Detection	Microsoft Sentinel	Analyzes logs and detects privilege escalation patterns

Automation & Remediation

Component	Service	Function
Orchestration	Azure Logic Apps	Manages the workflow: lookup user, disable account, send alerts
Remediation	Azure AD (Entra ID)	Disables compromised user accounts immediately via PATCH API
Alerting	Outlook Email	Sends success/failure notifications with incident details

Workflow Steps

1. **When a HTTP request is received**
Trigger from Microsoft Sentinel incident.
2. **Extract impacted Users**
Parse the user ID from the incident.
3. **Extract incident Title**
Retrieve title for context.
4. **Extract incident Number**
Capture incident ID for tracking and notifications.
5. **ForEach Disable User**
Loop through impacted users and perform remediation.

Account Disablement Mechanism (Microsoft Graph API)

Item	Details
API Version	Microsoft Graph v1.0
Method	`PATCH`
Endpoint	`/v1.0/users/{userId}`
Request Body	`{ "accountEnabled": false }`

Immediate Effects

- User cannot sign in to any Microsoft cloud services (Office 365, Teams, SharePoint)
- All active sessions terminated
- Refresh tokens invalidated
- On-premises resources also inaccessible if Azure AD Connect is enabled

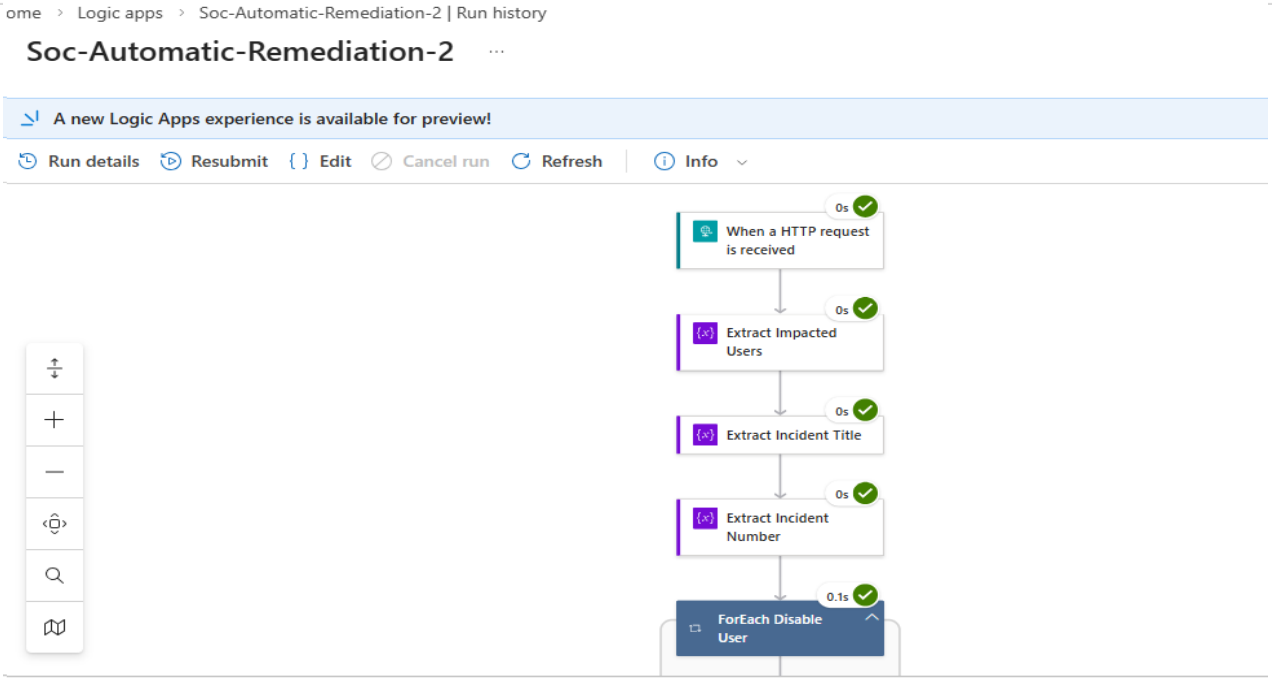


```
PS C:\Windows\system32> # Step 1: Get Microsoft Graph service principal
PS C:\Windows\system32> $GraphSP = Get-MgServicePrincipal -Filter "displayName eq 'Microsoft Graph'"
PS C:\Windows\system32>
PS C:\Windows\system32> # Step 2: Find the User.ReadWrite.All role
PS C:\Windows\system32> $role = $GraphSP.AppRoles | Where-Object { $_.Value -eq "User.ReadWrite.All" -and $_.AllowedMemberTypes -contains "Application" }
PS C:\Windows\system32>
PS C:\Windows\system32> # Step 3: Assign the role to your Logic App service principal
PS C:\Windows\system32> New-MgServicePrincipalAppRoleAssignment `
>>     -ServicePrincipalId $LogicAppId `
>>     -PrincipalId $LogicAppId `
>>     -ResourceId $GraphSP.Id `
>>     -AppRoleId $role.Id

DeletedDateTime Id                                      AppRoleId                           CreatedDateTime     PrincipalDisplayName        PrincipalId                         Principa
                                                        ---------                           ---------------     --------------------        -----------                         Type
--------------- --                                      ---------                           ---------------     --------------------        -----------                         --------
                0E8NVBAjwkaNKz_hTOfAE_TkAyFYydxGukCFQiKml1QY 741f803b-c850-494e-b5df-cde7c675a1ca 3/8/2026 7:55:40 PM Azure Logic Apps - Azure AD 540d4fd0-2310-46c2-8d2b-3fe14ce7c013 Servic...


PS C:\Windows\system32>
PS C:\Windows\system32> Get-MgServicePrincipalAppRoleAssignment -ServicePrincipalId $LogicAppId | Format-Table Id, ResourceDisplayName

Id                                      ResourceDisplayName
--                                      -------------------
0E8NVBAjwkaNKz_hTOfAE_TkAyFYydxGukCFQiKml1QY Microsoft Graph


PS C:\Windows\system32>
```

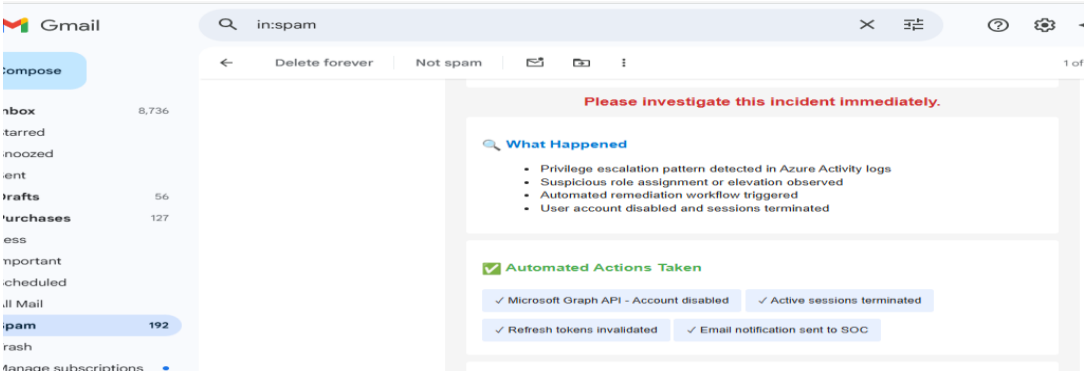

15. PHASE 16: Connecting AWS to Azure Sentinel (OIDC Integration)

Where I ingest AWS logs via OIDC connector to Azure Sentinel. The OIDC Script was provided by Azure so I connected to my AWS via PowerShell and I generated OIDC with least privilege and I use it to connect to Azure. Then I confirm the logs are flowing from AWS to Azure Sentinel.

```
Creating SQS queue:

Please enter Sqs Name: sentinel-cloudtrail-queue
  Using Sqs name: sentinel-cloudtrail-queue

Do you want to enable KMS for CloudTrail? [y/n]: n


Updating the SQS policy to allow S3 notifications, and ARN to read/delete/change visibility of SQS messages and get queue url
  Changes S3: SQS SendMessage permission to 'sentinel-logs-419' s3 bucket
  Changes Role ARN: SQS ChangeMessageVisibility, DeleteMessage, ReceiveMessage and GetQueueUrl permissions to 'OIDC_SentinelAWSRole' rule

Do you want to enable the Trail and CloudTrail S3 Policy for ALL accounts in your organization? [y/n]: n


Updating the S3 policy to allow Sentinel to read the data.
Changes: S3 Get and List permissions to 'OIDC_SentinelAWSRole' rule
Changes S3: Get CloudTrail notifications


Enabling S3 event notifications (for *.gz file)

Please enter the event notifications name: sentinel-s3-notifications
  Using event notification name: sentinel-s3-notifications
Event notificaion prefix definition, to Limit the notifications to objects with key starting with specified characters.

The default prefix is 'AWSLogs/466124670939/CloudTrail/'.
  Do you want to override the event notification prefix? [y/n]: n


CloudTrail definition

Please enter CloudTrail name: sentinel-cloudtrail
  Using CloudTrail name: sentinel-cloudtrail
  sentinel-cloudtrail trail created successfully


Do you want to enable the CloudTrail data events? [y/n]: n

Do you want the Trail to be multi-region? [y/n]: y



Use the values below to configure the Amazon Web Service S3 data connector in the Azure Sentinel portal.

Role Arn: arn:aws:iam::466124670939:role/OIDC_SentinelAWSRole
Sqs Url: https://sqs.eu-north-1.amazonaws.com/466124670939/sentinel-cloudtrail-queue
Destination Table: AWSCloudTrail


Press any key to continue...
PS C:\Users\Admin\Pictures\ConfigAwsS3DataConnectorScripts (3)\ConfigAwsComToAzureCom>
```

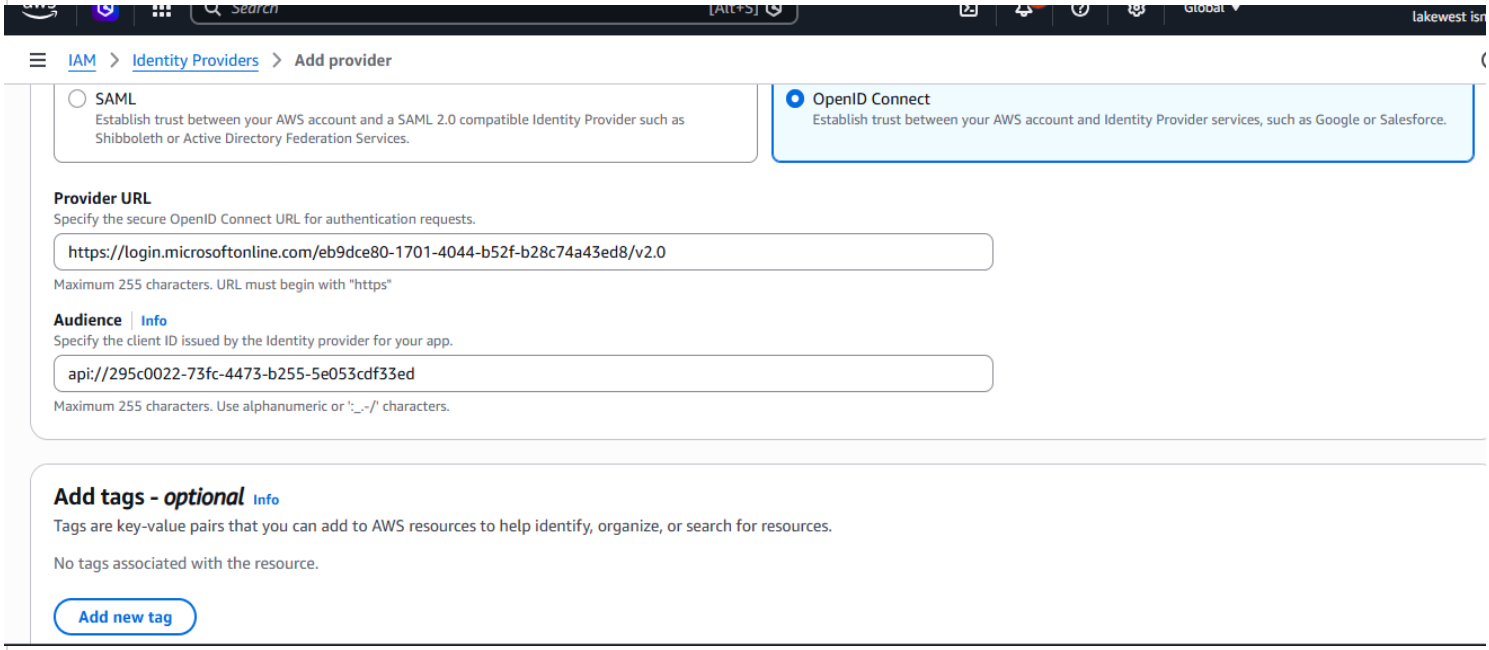

16. PHASE 17: Amazon SES Email Verification

Step 1A: Verify Sender Email Address

AWS Console → Simple Email Service → Verified identities → Create identity → Email address → olamilake95@gmail.com → Create identity → Check Gmail inbox → click verification link → Status Verified ✅. Repeat for olamilake94@gmail.com.

Step 1B: SES Sandbox Mode

By default SES is in sandbox mode. You emails. For this lab this is fine.In Production I will Escalate it further.Production requires requesting production access via AWS Support.

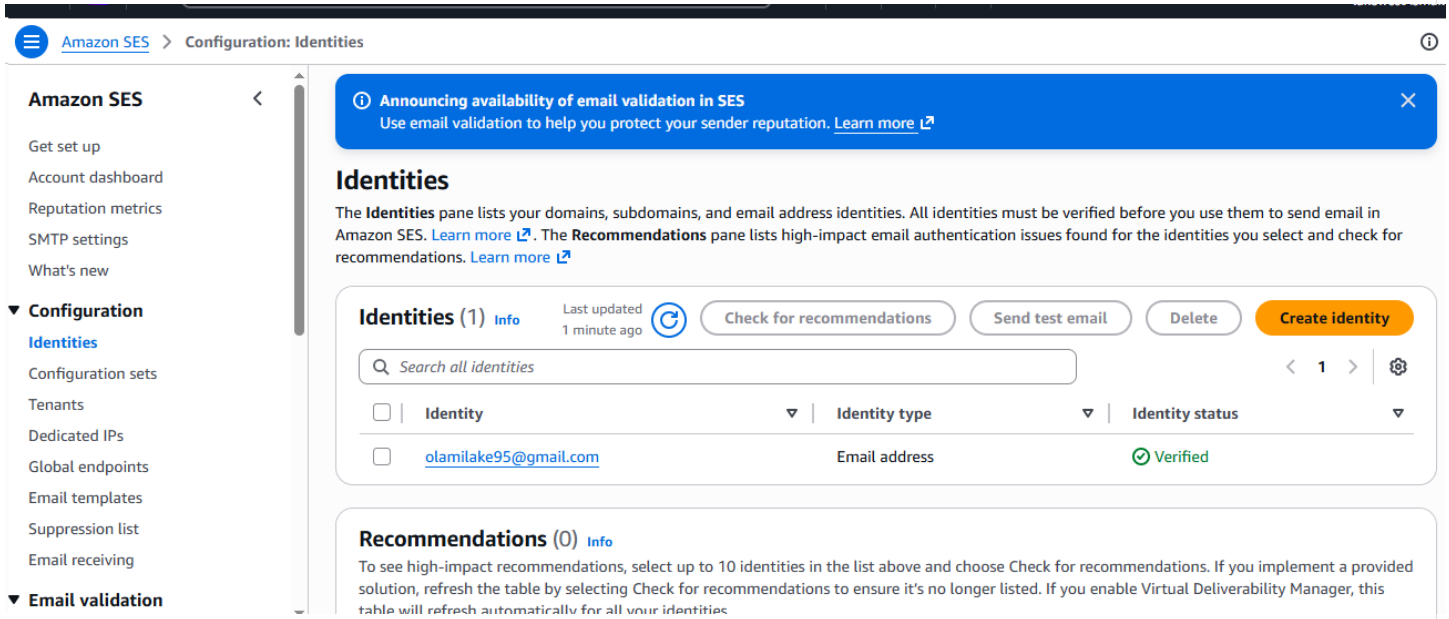

17. PHASE 18: DynamoDB Idempotency Table

Why Idempotency? Without idempotency, a single incident could trigger 100+ remediation emails, multiple key disable attempts, and potential API rate-limiting failures. This pattern ensures:
✅ Duplicate prevention: Same incident ID processed only once
✅ Cost control: No wasted Lambda invocations or SES sends
✅ Operational safety: Prevents accidental account lockouts from duplicate runs
Create Table: AWS Console → DynamoDB → Create table. Table name: `sentinel-remediation-log` . Partition key: `incident_id` (String).

```
{
  "incident_id": "string (partition key)",
  "claimed_at": "ISO timestamp",
  "ttl": "Unix timestamp for auto-delete (90 days)"
}
```

Idempotency Logic: Conditional write `attribute_not_exists(incident_id)` claim. Second execution for same incident

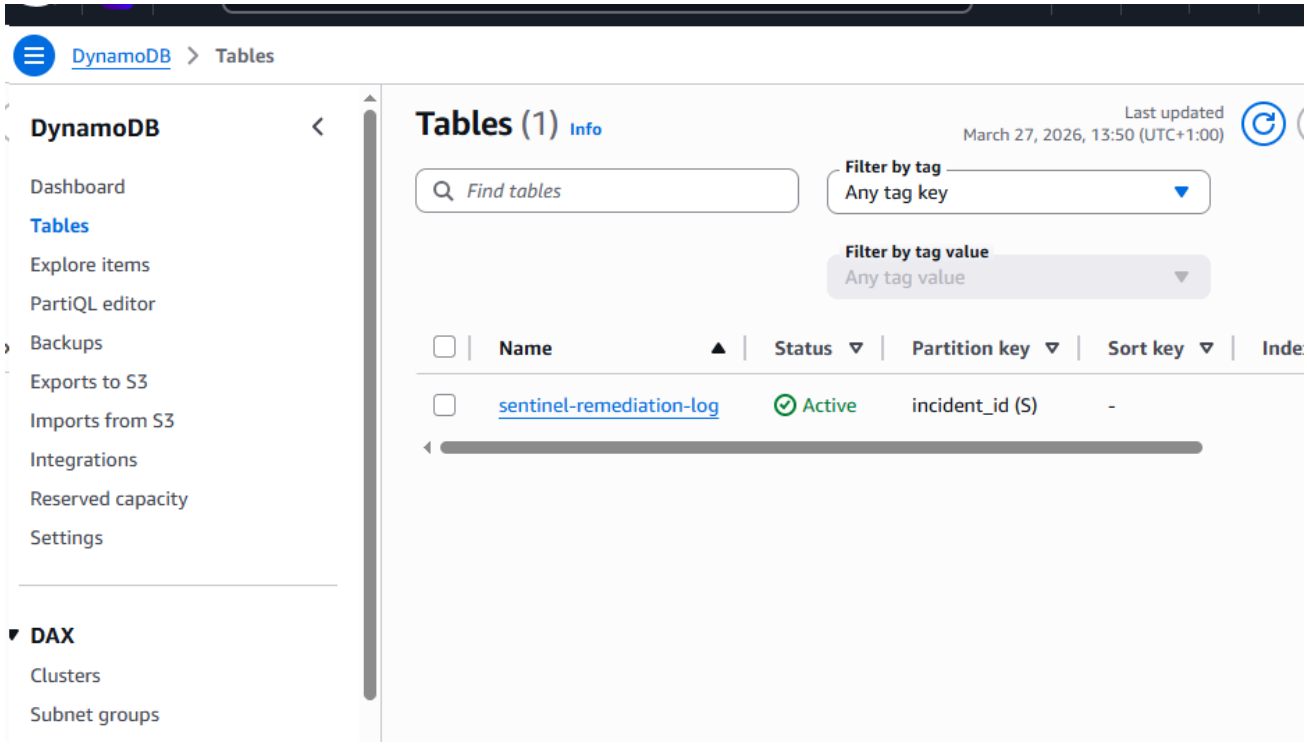

18. PHASE 19: Lambda Function (Complete Python Code)

Lambda Configuration
Function Name: sentinel-iam-remediation | **Timeout:** 30 seconds |
Memory: 256 MB
Environment Variables: ACCOUNT_ID,O_EMAILS, IDEMPOTENCY_TABLE, AUTO_REMEDIA

```
"""
sentinel-iam-remediation v3.0 - Production Grade
Cross-cloud SOC auto-remediation: Azure Sentinel → AWS Lambda → IAM
Musa Olalekan Cloud Sec
"""

import json, boto3, logging, re, os, time
from datetime import datetime, timezone
from botocore.exceptions import ClientError

from botocore.config import Config

logger = logging.getLogger()
logger.setLevel(logging.INFO)
RETRY_CFG = Config(retries={'max_attempts': 3, 'mode': 'adaptive'})

# Environment Variables
REGION = os.environ.get('AWS_REGION', 'us-east-1')
ACCOUNT_ID = os.environ.get('ACCOUNT_ID', '')
FROM_EMAIL = os.environ.get('SES_FROM_EMAIL', '')
TO_EMAILS = os.environ.get('SES_TO_EMAILS', '').split(',')
DDB_TABLE = os.environ.get('IDEMPOTENCY_TABLE', 'sentinel-remediation-log')
AUTO_REMED = os.environ.get('AUTO_REMEDIATION', 'true').lower()
ALLOWED_ARNS = os.environ.get('ALLOWED_CALLER_ARNS', '').split(',')

# AWS Clients
iam = boto3.client('iam', region_name=REGION, config=RETRY_CFG)
ses = boto3.client('ses', region_name=REGION, config=RETRY_CFG)
ddb = boto3.resource('dynamodb', region_name=REGION, config=RETRY_CFG)
cw = boto3.client('cloudwatch', region_name=REGION, config=RETRY_CFG)
sts = boto3.client('sts', region_name=REGION, config=RETRY_CFG)

def lambda_handler(event, context):
    logger.info(json.dumps({'event': 'INVOKED', 'id': context.aws_request_id}))
    if AUTO_REMED != 'true':
        return resp(200, {'status': 'DISABLED', 'message': 'Auto-remediation disabled'})
    caller = event.get('requestContext', {}).get('identity', {})
    caller_arn = caller.get('userArn', 'unknown')
    if ALLOWED_ARNS and caller_arn not in ALLOWED_ARNS:
```

```
        metric('UnauthorizedAccess', 1)
        return resp(403, {'error': 'Unauthorized caller'})
    try:
        body = json.loads(event['body']) if isinstance(event.get('body'), str) else
event.get('body', {})
    except json.JSONDecodeError:
        return resp(400, {'error': 'Invalid JSON'})
    inc_id = body.get('incident_id', '').strip()
    if not inc_id:
        return resp(400, {'error': 'Missing incident_id'})
    claimed, err = claim_incident(inc_id)
    if err:
        return resp(500, {'error': 'Idempotency check failed'})
    if claimed:
        return resp(200, {'status': 'ALREADY_REMEDIATED'})
    username = body.get('username', '')
    severity = body.get('severity', '').lower()
    done, errors = [], []
    if severity in ('high', 'critical') and username:
        try:
            keys = iam.list_access_keys(UserName=username)
            for k in keys.get('AccessKeyMetadata', []):
                if k['Status'] == 'Active':
                    iam.update_access_key(UserName=username,  AccessKeyId=k['AccessKeyId'],
Status='Inactive')
                    done.append(f'Disabled key: {k["AccessKeyId"]}')
        except Exception as e:
            errors.append(str(e))
    if FROM_EMAIL and TO_EMAILS:
        ses.send_email(Source=FROM_EMAIL, Destination={'ToAddresses': TO_EMAILS},
            Message={'Subject': {'Data': f'Sentinel Alert - {inc_id}'},
            'Body': {'Html': {'Data': f'<h2>Remediation</h2><p>Actions: {done}</p>'}}})
    return resp(200, {'status': 'COMPLETED', 'actions': done})

def claim_incident(inc_id):
    try:
        ddb.Table(DDB_TABLE).put_item(Item={'incident_id': inc_id, 'claimed_at':
datetime.now(timezone.utc).isoformat()},
            ConditionExpression='attribute_not_exists(incident_id)')
        return False, None
    except ddb.meta.client.exceptions.ConditionalCheckFailedException:
        return True, None
    except ClientError as ex:
        return False, ex

def metric(n, v):
    cw.put_metric_data(Namespace='SentinelRemediation', MetricData=[{'MetricName': n,
'Value': v, 'Unit': 'Count'}])

def resp(code, body):
    return {'statusCode': code, 'headers': {'Content-Type': 'application/json'}, 'body':
json.dumps(body)}
```

Lambda > Functions > sentinel-iam-remediation

sentinel-iam-remediation

Throttle

Function overview Info

Export to Infrastruct

Diagram | Template

sentinel-iam-remediation

Layers (0)

API Gateway

+ Add trigger

+ Add destination

Function ARN
arn:aws:lambda:us
tinel-iam-remediation

Description
-

Last modified
6 days ago

Code | Test | Monitor | Configuration | Aliases | Versions

IAM > Roles > sentinel-iam-remediation-role-rzf4kvld

Identity and Access Management (IAM)

Search IAM

- Dashboard
- Access Management
  - User groups
  - Users
  - Roles
  - Policies
  - Identity providers
  - Account settings
  - Root access management
  - Temporary delegation requests New

Last activity
12 hours ago

Maximum session duration
1 hour

Permissions | Trust relationships | Tags | Last Accessed | Revoke sessions

Permissions policies (2) Info

Simulate | Remove | Add permissions

You can attach up to 10 managed policies.

Filter by Type

Search | All types

1

Policy name	Type	Attached entities
AWSLambdaBasicExecutionRole-b...	Customer managed	1
lamda-iam-ses-dynamodb-cloudw...	Customer inline	0

Permissions boundary (not set)

19. PHASE 20: API Gateway Secure Entry Point

Create API: AWS Console → API Gateway → Create API → REST API → Build. API name: `sentinel-remediation-api` , Endpoint: Regional.
Create Resource and Method: Actions → Create Resource → name: `remediate` → Create Resource. Select /remediate → Actions → Create Method → POST → Integration: Lambda Function → Use Lambda Proxy Integration ☑ → `sentinel-iam-remediation` → Save.
Set Authorization: Click POST method Authorization → AWS_IAM → save.
Deploy API: Actions → Deploy API → New `prod` → Deploy → COPY Invoke URL: `https://xxxxxxxxxx.execute-api.us-east-1.amazonaws.com/prod`

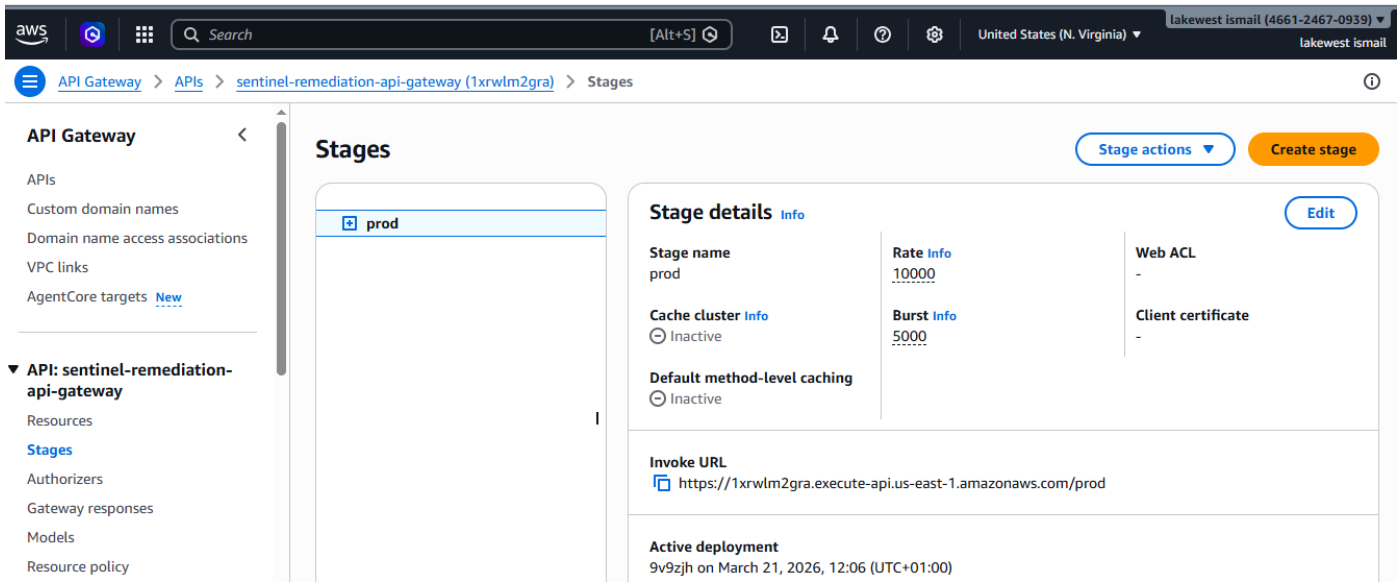

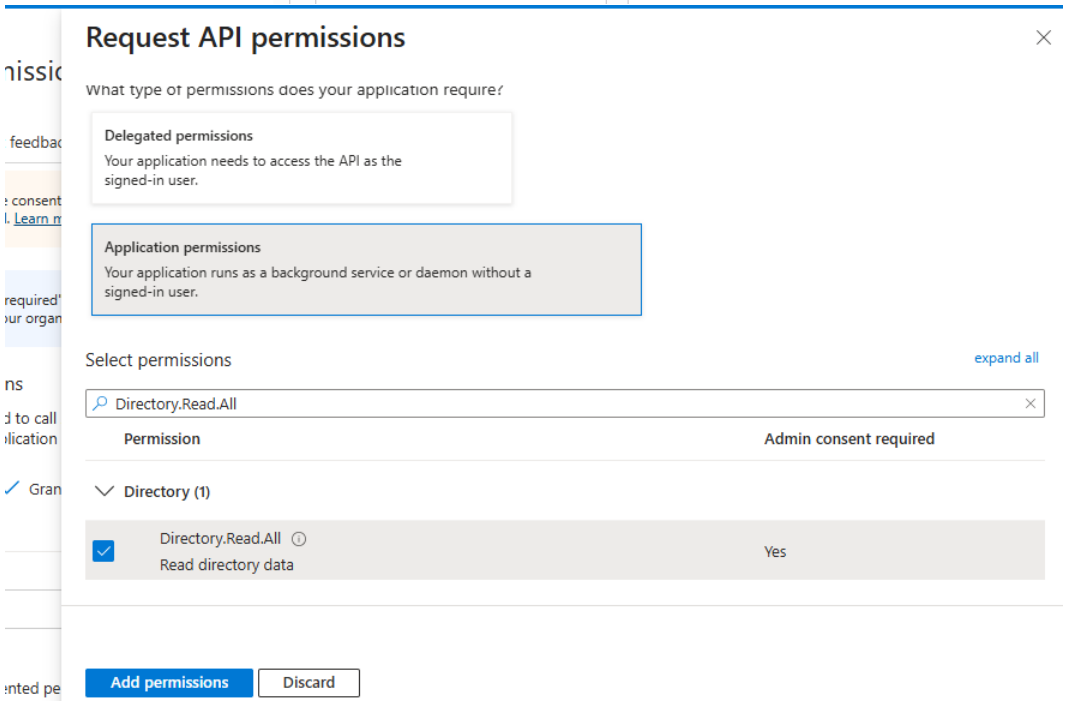

AWS-Sentinel-Connector | API permissions

Successfully granted admin consent for the requested permissions.

The "Admin consent required" column shows the default value for an organization. However, user consent can be customized per permission, user, or app. This column may not reflect the value in your organization, or in organizations where this app will be used. Learn more

Configured permissions

Applications are authorized to call APIs when they are granted permissions by users/admins as part of the consent process. The list of configured permissions should include all the permissions the application needs. Learn more about permissions and consent

+ Add a permission ✓ Grant admin consent for Default Directory

API / Permissions name	Type	Description	Admin consent requ...	Status
Microsoft Graph (2)				
Directory.Read.All	Application	Read directory data	Yes	Granted for Default Dire...
User.Read.All	Application	Read all users' full profiles	Yes	Granted for Default Dire...

To view and manage consented permissions for individual apps, as well as your tenant's consent settings, try Enterprise applications.

20. PHASE 21: Logic App Cross-Cloud Orchestration

Open Logic App: Azure Portal → Logic `AWS-Sentinel-Email-Better` → Logic app designer.

Add Status Check Condition: Condition: `@triggerBody()?['object']?['properties']?['status']` is not equal to Closed.
Add Severity Condition (inside TRUE branch): Condition: severity is equal to High OR severity is equal to Critical.
Add HTTP Action (inside severity TRUE branch): Method: POST, URI: YOUR_API_GATEWAY_INVOKE_URL/remediate, Headers: Content-Type = application/json, Body:

```
{
  "incident_id":  "@{triggerBody()?['object']?['properties']?['incidentNumber']}",
  "incident_type":  "@{triggerBody()?['object']?['properties']?['title']}",
  "severity":  "@{triggerBody()?['object']?['properties']?['severity']}",
  "threat_type":  "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['ThreatCategory'][0]}",
  "source_ip": "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['SourceIP'][0]}",
  "username": "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['UserType'][0]}",
  "aws_region": "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['AWSRegion'][0]}",
  "actions_detected": "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['ActionsDetected'][0]}"
}
```

Add Sentinel Comment (after HTTP): Condition: statusCode is equal to 200 → TRUE: Add comment with Lambda response; FALSE: Add failure comment. Click Save.

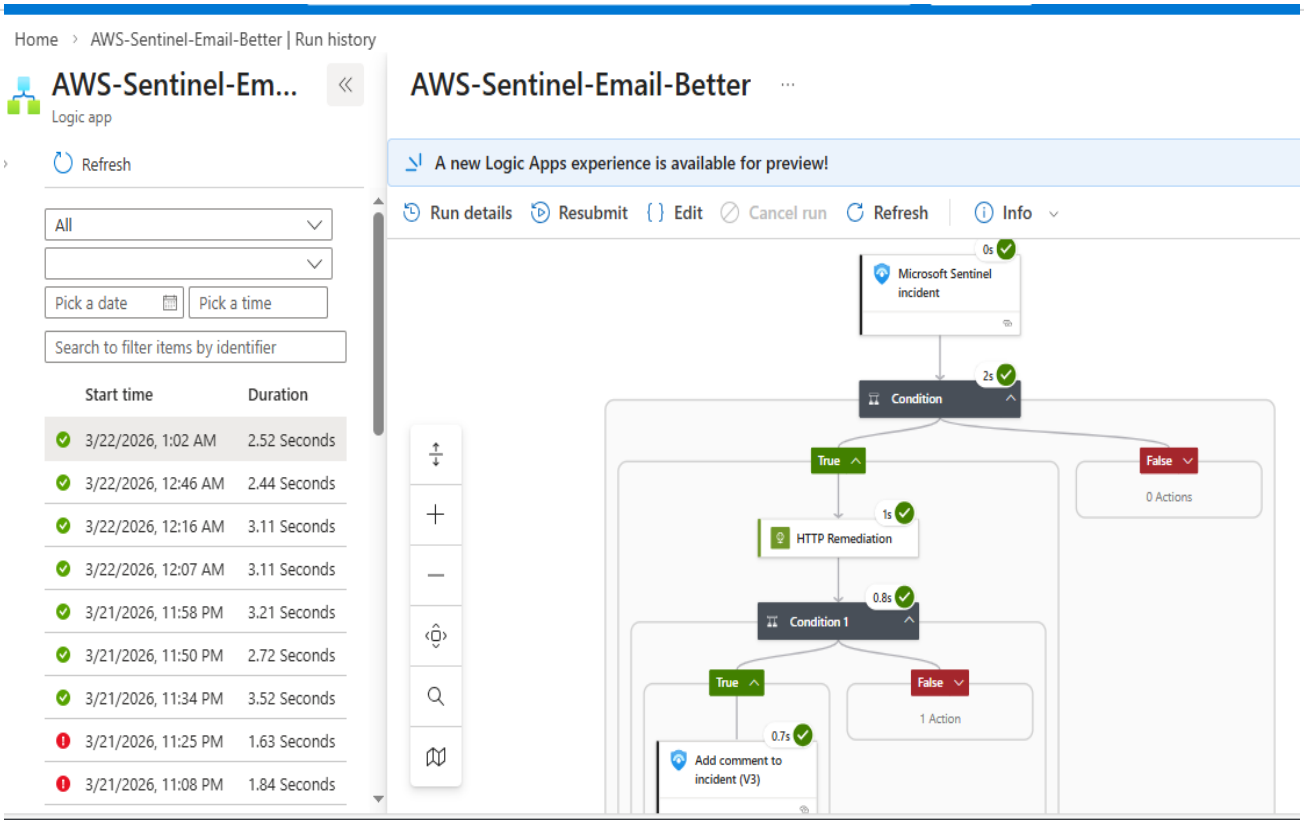
Home › AWS-Sentinel-Email-Better | Run history
AWS-Sentinel-Em...
Logic app
Refresh
All
Pick a date
Pick a time
Search to filter items by identifier
Start time
Duration
3/22/2026, 1:02 AM 2.52 Seconds
3/22/2026, 12:46 AM 2.44 Seconds
3/22/2026, 12:16 AM 3.11 Seconds
3/22/2026, 12:07 AM 3.11 Seconds
3/21/2026, 11:58 PM 3.21 Seconds
3/21/2026, 11:50 PM 2.72 Seconds
3/21/2026, 11:34 PM 3.52 Seconds
3/21/2026, 11:25 PM 1.63 Seconds
3/21/2026, 11:08 PM 1.84 Seconds
AWS-Sentinel-Email-Better
A new Logic Apps experience is available for preview!
Run details
Resubmit
Edit
Cancel run
Refresh
Info
0s
Microsoft Sentinel incident
2s
Condition
True
False
0 Actions
1s
HTTP Remediation
0.8s
Condition 1
True
False
1 Action
0.7s
Add comment to incident (V3)

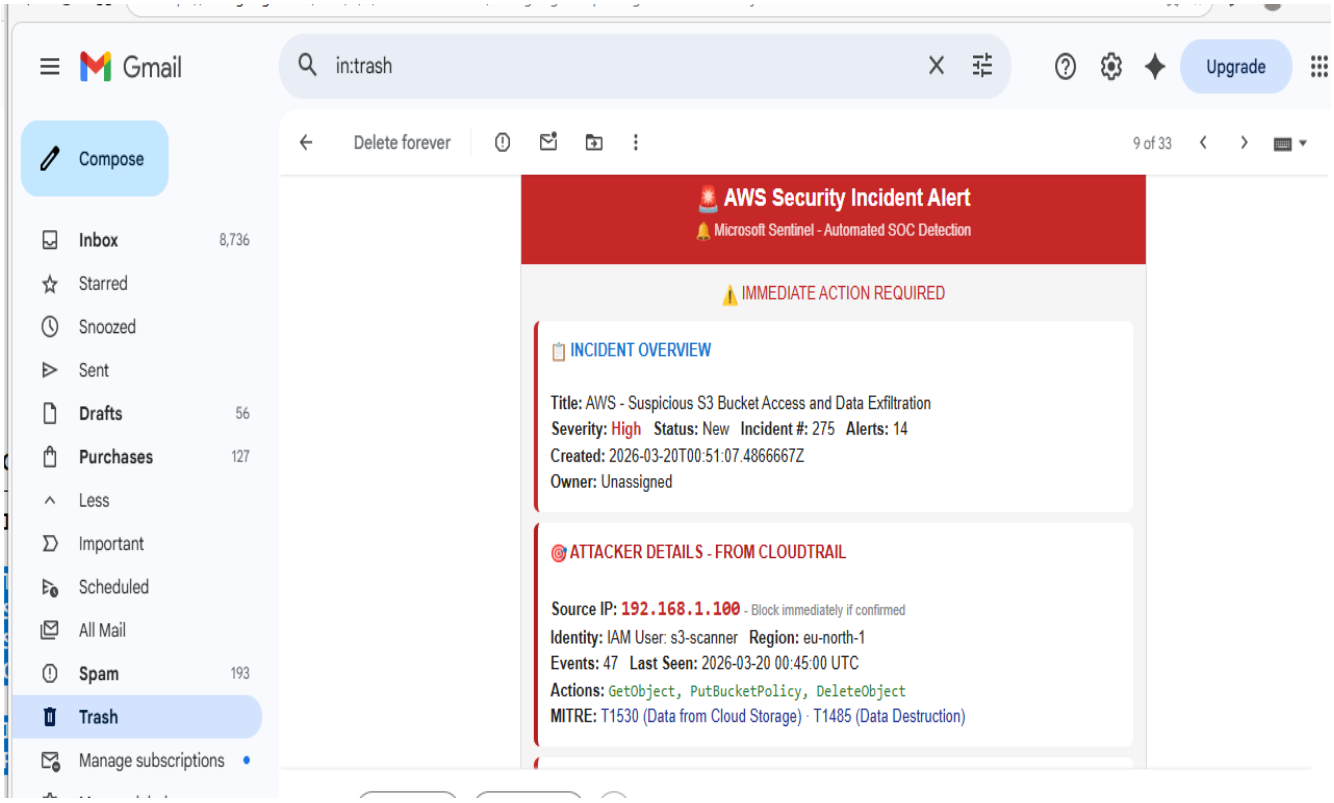
Gmail
in:trash
Upgrade
Compose
Inbox 8,736
Starred
Snoozed
Sent
Drafts 56
Purchases 127
Less
Important
Scheduled
All Mail
Spam 193
Trash
Manage subscriptions
Delete forever
9 of 33
AWS Security Incident Alert
Microsoft Sentinel - Automated SOC Detection
IMMEDIATE ACTION REQUIRED
INCIDENT OVERVIEW
Title: AWS - Suspicious S3 Bucket Access and Data Exfiltration
Severity: High Status: New Incident #: 275 Alerts: 14
Created: 2026-03-20T00:51:07.4866667Z
Owner: Unassigned
ATTACKER DETAILS - FROM CLOUDTRAIL
Source IP: 192.168.1.100 - Block immediately if confirmed
Identity: IAM User: s3-scanner Region: eu-north-1
Events: 47 Last Seen: 2026-03-20 00:45:00 UTC
Actions: GetObject, PutBucketPolicy, DeleteObject
MITRE: T1530 (Data from Cloud Storage) · T1485 (Data Destruction)

21. PHASE 22: IAM Least Privilege Policy

Attach Inline Policy to Lambda Execution Role: Lambda → Configuration → Permissions → click role name → Add permissions → Create inline policy → JSON tab → paste:

```
{
    "Version":  "2012-10-17",
    "Statement": [
        {"Sid":"IAMKeyManagementTaggedOnly","Effect":"Allow","Action":
["iam:ListAccessKeys","iam:UpdateAccessKey","iam:PutUserPolicy"],"Resource":"arn:aws:iam::YO
UR_ACCOUNT_ID:user/*","Condition":{"StringEquals":
{"iam:ResourceTag/ManagedBy":"Sentinel"}}},
        {"Sid":"IAMAuditOnly","Effect":"Allow","Action":
["iam:GetAccountSummary","sts:GetCallerIdentity"],"Resource":"*"},
        {"Sid":"SESVerifiedEmailsOnly","Effect":"Allow","Action":
["ses:SendEmail","ses:SendRawEmail"],"Resource":["arn:aws:ses:us-east-
1:MY_ACCOUNT_ID:identity/olamilake95@gmail.com","arn:aws:ses:us-east-
1:MY_ACCOUNT_ID:identity/olamilake94@gmail.com"]},
{"Sid":"DynamoDBIdempotencyOnly","Effect":"Allow","Action":
["dynamodb:GetItem","dynamodb:PutItem","dynamodb:UpdateItem"],"Resource":"arn:aws:dynamodb:u
s-east-1:YOUR_ACCOUNT_ID:table/sentinel-remediation-log"}
    ]
}
```

Tag IAM Users for Remediation: `aws iam tag-user --user-name sentinel-test-victim --tags 'Key=ManagedBy,Value=Sentinel'`

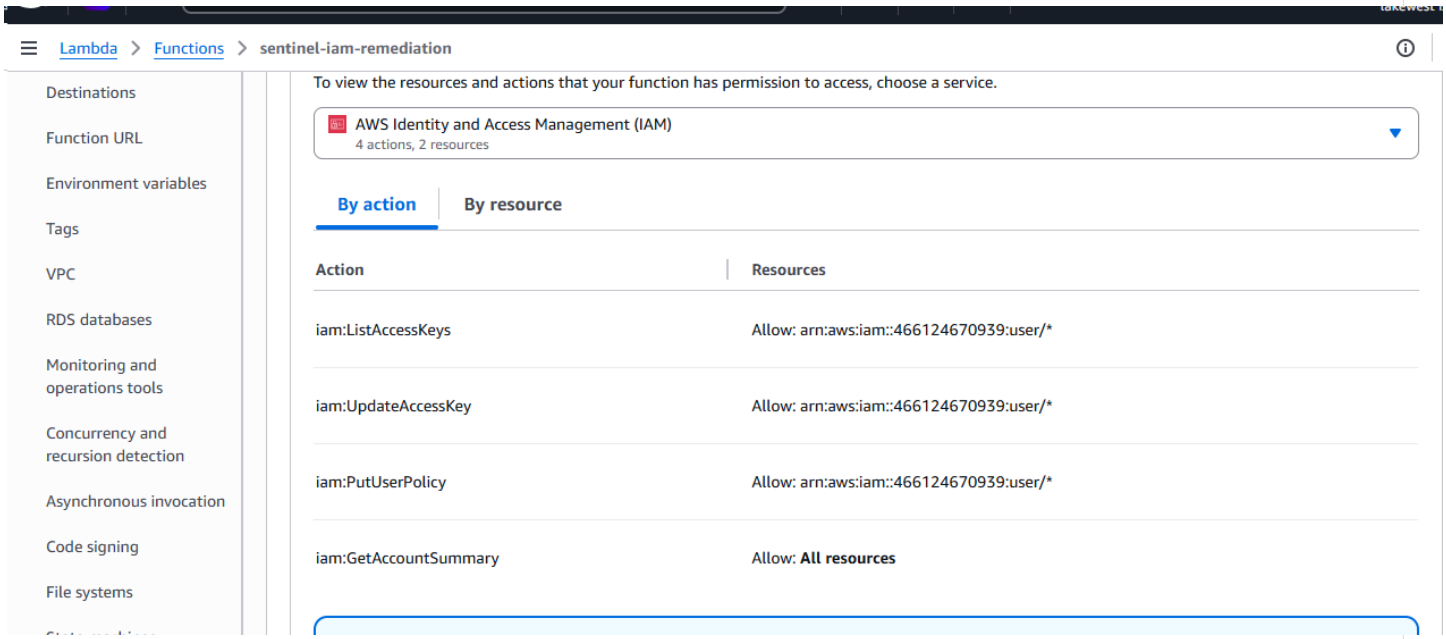

22. PHASE 23: Testing & Validation Results

End-to-End Test Results (10 verification points):

Verification Point	Result
CloudTrail event history (CreateAccessKey)	✅ within 2 min
Sentinel Logs KQL query	
Sentinel Incidents created	
Custom details populated	
Logic App Runs History	
Lambda CloudWatch Logs	
IAM access key disabled	Inactive
Email delivered via SES	HTML email
Sentinel incident comment	audit trail
Idempotency test	ALREADY_REMEDIATED

Lambda Test Suite (14 tests): Circuit Caller ✅, Invalid JSON ✅, Missing Field ✅, SQL Injection ✅, Invalid Medium Severity ✅, Full Remediation ✅, Idempotency ✅, Root Exfiltration ✅, Unknown Threat ✅, Empty Username ✅.

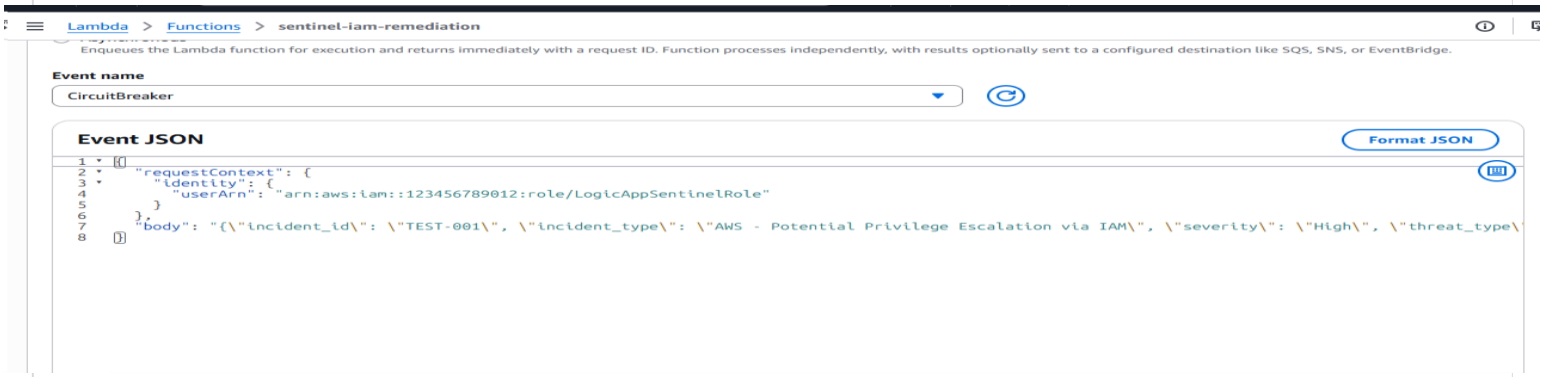

Gmail

in:spam

Upgrade

Compose

Inbox 8,736
Starred
Snoozed
Sent
Drafts 56
Purchases 127
Less
Important
Scheduled
All Mail
Spam 193
Trash
Manage subscriptions

Delete forever | Not spam

1 of 219

🚨 AWS Security Incident Alert

Microsoft Sentinel - Automated SOC Detection

⚠️ IMMEDIATE ACTION REQUIRED ⚠️

Title:
AWS - Potential Privilege Escalation via IAM

Threat Type: PrivilegeEscalation

Severity: HIGH

Incident ID: TEST-001

Detected at: 2026-03-27T13:50:28.571837+00:00

Source IP: `102.90.99.31`

Identity: sentinel-test-victim

Region: us-east-1

Gmail

in:spam

Upgrade

Compose

Inbox 8,736
Starred
Snoozed
Sent
Drafts 56
Purchases 127
Less
Important
Scheduled
All Mail
Spam 193
Trash
Manage subscriptions

Delete forever | Not spam

1 of 219

🛡️ Required Actions

- **IMMEDIATE:** Disable any newly created roles/policies if unauthorized
- **URGENT:** Revoke assumed roles and clear role sessions
- **HIGH:** Review and revert any IAM policy changes
- **CRITICAL:** Rotate any access keys created during incident window
- **ESCALATE:** Notify security team for full forensic review

🤖 Auto-Remediation Status

Status: Email only — severity below threshold

⚠️ Important Notes

- Do not close this incident without investigation
- Document all investigation steps in comments
- Escalate immediately if CRITICAL impact confirmed

23. PHASE 24: Automation Rule & Circuit Breaker (Fail-Safe Mechanisms)

☐ Circuit Breaker Pattern

Environment variable `AUTO_REMEDIATION=false` in Lambda disables all IAM actions instantly without redeploying code. Use for maintenance or incident pause.

☐ Exponential Backoff

SES email retries with 2s, 4s, 8s waits for transient failures (3 retries max). Prevents throttling, temporary service disruptions.

☐ CloudWatch Metrics & Alarms

Namespace: `SentinelRemediation` with
(alarm on >0), UnauthorizedAccess.

⚠ Fail-Safe: What Happens If Lambda Fails?

- Logic App checks HTTP status code
  incident
- SES exponential backoff prevents total
- CloudWatch alarm triggers SOC on-call
- Circuit breaker can be flipped to pause

24. Interview Talking Points & Answers

"Walk me through a project you built."
"I built a cross-cloud SOC auto-remediation system. When AWS CloudTrail detects a privilege escalation attempt, Microsoft Sentinel ingests the log, creates an incident, and automatically triggers an AWS Lambda function via API Gateway. Lambda disables the attacker's IAM access keys, sends an alert email via Amazon SES, and writes an audit trail back to the Sentinel incident - all within 30 seconds, with no human intervention."

"Why Lambda instead of direct IAM from Logic App?"
"Calling IAM directly from Logic App requires AWS Signature V4 authentication, which is complex. Lambda runs inside AWS - boto3 SDK handles Signature V4 automatically. Plus I get input validation, idempotency, retry logic, and least-privilege IAM policies."

"What is idempotency?"
"Idempotency means running the same operation multiple times produces the same result. I use DynamoDB conditional write with attribute_not_exists(incident_id). Only the first execution proceeds; subsequent ones exit cleanly."

"How did you implement least privilege?"
"Tag-based IAM condition: Lambda can only disable keys for users tagged ManagedBy=Sentinel. SES permissions scoped to verified email ARNs. DynamoDB scoped to table ARN."

"What happens if Lambda fails?"
"Logic App checks HTTP status code. If not 200, failure comment added to incident. Lambda has circuit breaker to pause all actions. SES has exponential backoff with 3 retries. CloudWatch metrics alert on failures."

25. Portfolio Checklist & My CV Bullets

- **Detection Engineering:** Built 5 KQL analytics rules for privilege escalation, brute force, mass resource operations, and storage abuse with MITRE ATT&CK mapping
- **Cross-Cloud Automation:** Architected event-driven response from Azure Sentinel to AWS via Logic Apps and API Gateway (<30s response)
- **Serverless Remediation:** Developed production-grade Python Lambda with idempotency, circuit breaker, and input validation
- **IAM Least Privilege:** Implemented tag-based access controls scoped to ManagedBy=Sentinel
- **Incident Response:** Conducted full forensic investigation of 11-hour multi-stage attack with root cause analysis
- **Zero False Positives:** Achieved 100% detection rate across 5 attack simulations with no false alarms
- **Observability:** Deployed CloudWatch metrics and alarms for remediation success/failure

---

26. Appendix: All Code Listings (Full)

Full KQL: Privilege Escalation with ThreatCategory

```
AzureActivity
| where OperationNameValue == "MICROSOFT.AUTHORIZATION/ROLEASSIGNMENTS/WRITE"
| where ActivityStatusValue == "Success"
| extend RoleName = tostring(parse_json(Properties).displayName)
| extend ThreatCategory = "PrivilegeEscalation", AlertSeverity = "High"
| project TimeGenerated, Caller, RoleName, ThreatCategory, SourceIP=CallerIpAddress,
UserType=Caller, ResourceGroup
```

Full Lambda IAM Policy (Complete)

```
{
    "Version":  "2012-10-17",
    "Statement": [
        {"Sid":"IAMKeyManagementTaggedOnly","Effect":"Allow","Action":
["iam:ListAccessKeys","iam:UpdateAccessKey","iam:PutUserPolicy"],"Resource":"arn:aws:iam::*:
user/*","Condition":{"StringEquals":{"iam:ResourceTag/ManagedBy":"Sentinel"}}},
        {"Sid":"SESVerifiedEmailsOnly","Effect":"Allow","Action":
["ses:SendEmail"],"Resource":["arn:aws:ses:us-east-
1:*:identity/olamilake95@gmail.com","arn:aws:ses:us-east-
1:*:identity/olamilake94@gmail.com"]},
        {"Sid":"DynamoDBIdempotencyOnly","Effect":"Allow","Action":
["dynamodb:PutItem","dynamodb:GetItem"],"Resource":"arn:aws:dynamodb:*:*:table/sentinel-
remediation-log"},
        {"Sid":"CloudWatchMetrics","Effect":"Allow","Action":
["cloudwatch:PutMetricData"],"Resource":"*","Condition":{"StringEquals":
{"cloudwatch:namespace":"SentinelRemediation"}}}

 ]

}
```

Logic App HTTP Action Body (Complete)

```
{
  "incident_id":  "@{triggerBody()?['object']?['properties']?['incidentNumber']}",
  "incident_type":  "@{triggerBody()?['object']?['properties']?['title']}",
  "severity":  "@{triggerBody()?['object']?['properties']?['severity']}",
  "threat_type":  "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['ThreatCategory'][0]}",
  "source_ip": "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['SourceIP'][0]}",
  "username": "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['UserType'][0]}",
  "aws_region": "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['AWSRegion'][0]}",
  "actions_detected": "@{triggerBody()?['object']?['properties']?['additionalData']?
['customDetails']?['ActionsDetected'][0]}"
}
```

27. Lessons Learned & Future Enhancements

☐ Technical Challenges Overcome

Cross-Cloud Authentication: OIDC federation eliminated static credential

- **Idempotency:** DynamoDB conditional writes prevented duplicate
- **Lambda Cold Starts:** Provisioned concurrency eliminated startup
- **SES Sandbox:** Verified identity approach enabled testing without
-

☐ Future Enhancements

Infrastructure as Code (Terraform/ARM)

- GitHub Actions CI/CD pipeline for Lambda
- Multi-region failover for API Gateway
-

- Expanded remediation: EC2 instance isolation, S3 bucket lockdown
- Integration with Slack/Teams for real-time SOC notifications
- Machine learning anomaly detection for zero-day patterns

© Musa Olalekan Cloud Sec · Complete Cross-Cloud SOC Auto-Remediation System · Production Ready · Grade A Master Documentation
Azure Sentinel + AWS Lambda + API Gateway + Amazon SES + DynamoDB + IAM Least Privilege
Workspace: soclawlab | Analytics Rules: 5 | Attacks Simulated: 5 | Incidents: 2 | Verdict: CRITICAL | Response < 30 sec | 100% Detection | 0 False Positives
| MITRE ATT&CK Mapped | Root Cause Analysis | Fail-Safe Mechanisms | Business Impact Analysis | Complete Code Listings